

Using ring recruitment to power response-adaptive randomisation in vaccine trials during an epidemic

August 27, 2020

Abstract

A clinical trial to test a vaccine in the midst of an epidemic faces a number of unique challenges, including incidence varying in time and space, difficulty in establishing eligibility of potential participants, and a time pressure to find an efficacious agent in order to curtail an outbreak. Recently, novel vaccine trial designs were proposed to address these specific challenges, including targeted recruitment of people at imminent risk of infection, and determination of the trial's size based on the number of events seen rather than the number of people enrolled (Henao-Restrepo et al., 2015; WHO R&D Blueprint, 2020).

We develop a network-based simulation model to demonstrate and assess such trial designs, focusing on COVID-19 as an example. The simulations allow us to assess designs in terms of power and other operating characteristics of the trial, as well as the number of infections among participants.

Restricting our attention to two-arm trials and a binary outcome, we present and assess two design and analysis choices that are geared towards trials for vaccines for infectious diseases that spread person to person, such as COVID-19 and Ebola virus disease:

1. Recruiting contacts who are connected to known cases as a means to maximise number of events seen in a short time horizon.
2. Accounting for the possibility that individuals were infected before a vaccine had a chance to protect them by weighting the data in the analysis (which “retrospective exclusion”).

In the treatment-trial literature, adaptive designs are proposed to tailor trials in order to optimise some outcome, such as treatment successes or speed to conclusion. One such design, response-adaptive randomisation, is ill suited to vaccine trials as they are not useful for outcomes that take a long time to observe (Wason et al., 2019). Having identified a design that allows a vaccine trial to operate on a faster timescale, by (a) recruiting participants who are at imminent risk, and (b) weighting rather than excluding data points in order to maximise information retention, we overcome that obstacle, and thus enable response-adaptive randomisation in vaccine trials.

We conclude that, together, these design choices allow trials to be designed that meet the specific challenges of vaccine trials, including the objective of maximising population benefit. We identify some design and analysis elements that are demonstrably beneficial for certain infectious diseases. We present the use of response-adaptive randomisation and how simulation can be used to decide which method is appropriate.

Contents

1	Introduction	3
1.1	Background	3
1.2	Open challenges	3
1.3	Addressing the challenges	3
1.4	Outline	4
2	Simulation study	4
2.1	Trial design	4
2.2	Operating characteristics and metrics reported	5
3	Non-adaptive elements and analysis choices	5
3.1	Recruitment schedule: contact tracing	5
3.1.1	Methods	5
3.1.2	Results	6
3.1.3	Interpretation	6

47	3.2	Exclusion criterion implemented at analysis points	7
48	3.2.1	Methods	7
49	3.2.2	Results	7
50	3.2.3	Interpretation	7
51	3.3	Conclusions	8
52	4	Response-adaptive randomisation	9
53	4.1	Methods	9
54	4.1.1	Frequentist response-adaptive	9
55	4.1.2	Bayesian response-adaptive (Thompson sampling)	9
56	4.2	Results	12
57	5	Conclusion	13
58	5.1	Design and analysis choices	13
59	5.2	Adaptations	14
60	5.3	Follow-up time	14
61	5.4	Endpoint	15
62	5.5	Vaccine efficacy estimates	15
63	5.6	Future work	16
64	5.6.1	Relation to the epidemic	16
65	5.6.2	Formal comparison between response-adaptive randomisation and early stopping	16
66	A	The COVID-19 model	20
67	A.1	Network model	20
68	A.1.1	Definitions	21
69	A.1.2	Parametrisation	23
70	A.2	Disease and trial state transitions	24
71	A.2.1	Disease transmission rules	24
72	A.2.2	Trial rules	25
73	B	Analyses, and exclusion criterion implemented at analysis points	27
74	B.1	Analysis of raw data	27
75	B.2	Analysis using binary weighting	27
76	B.3	Analysis using continuous weighting	27
77	C	Supplementary material	29
78	C.1	Weighting example	29
79	C.2	Allocation probabilities	29
80	C.3	Time trend in incidence of infection	29
81	D	Trial size	31
82	E	Glossary by letter	36

1 Introduction

1.1 Background

Clinical trials for vaccines are currently in progress for SARS-CoV-2 (WHO R&D Blueprint, 2020), (<http://www.ox.ac.uk/news/2020-05-22-oxford-covid-19-vaccine-begin-phase-iii-human-trials>), and many vaccines are in development (www.who.int/publications/m/item/draft-landscape-of-covid-19-candidate-vaccines). So far the trial designs proposed have been two-arm, individually randomised placebo-controlled trials, with the exception of the WHO Blueprint, which allows for new arms to be added. In addition, it plans to carry out statistical tests for efficacy following the observation of a fixed number of cases, rather than following the enrolment of a fixed number of participants (WHO R&D Blueprint, 2020). Participants are (to be) recruited at random from within particular countries in response to incidence (WHO R&D Blueprint (2020), <https://www.ovg.ox.ac.uk/news/trial-of-oxford-covid-19-vaccine-starts-in-brazil>, <https://www.ovg.ox.ac.uk/news/trial-of-oxford-covid-19-vaccine-in-south-africa-begins>).

In the ongoing COVID-19 pandemic, many nations have employed contact tracing as part of their efforts to contain the spread of disease (ECDC, 2020), but to the best of our knowledge contact tracing systems have not been used as an embedded element of treatment or vaccine trials. In the outbreak of Ebola virus disease (EVD) from 2014 to 2016, one vaccine trial recruited participants according to infection events as they arose (Henao-Restrepo et al., 2017). Its ring recruitment strategy was based on contact tracing, and in this way identified those at imminent risk of infection following an international protocol (WHO Disease Surveillance and Response Programme Area Disease Prevention and Control Cluster, 2014).

A feature the two designs have in common is that they include adaptive design elements (WHO R&D Blueprint, 2020; Henao-Restrepo et al., 2017). An adaptive design is one in which changes to the design occur in a prespecified way as the trial progresses according to data that are accrued, with the aim of improving power, efficiency or participant benefit of clinical trials (Wason et al., 2019). The trial described in WHO R&D Blueprint (2020) allows for many planned adaptations, including what the candidate vaccines are, what the control arm is, where the trial is operating and recruiting participants, and early stopping. The Henao-Restrepo et al. (2017) trial planned for and conducted an interim analysis to allow for early stopping. They concluded the trial early, having found somewhat convincing evidence of the vaccine’s efficacy.

Among the suite of adaptive designs are interim analyses for response-adaptive randomisation (Kahn et al., 2020). An interim analysis is one which is planned for, in which the data accumulated so far are analysed in order to make pre-planned changes to the protocol. Adaptation to the randomisation ratio, in which the allocation ratio between experimental and control is adjusted, can be used to maximise power or patient benefit, or a combination of both (i.e. patient benefit subject to a power constraint).

1.2 Open challenges

It has been suggested that response-adaptive randomisation might improve a vaccine trial, in terms of efficiency or vaccination successes (Kahn et al., 2020). Such a design requires the endpoint, if binary, (or the event, if time to event) to be near in time (Wason et al., 2019). One way to achieve this is to recruit individuals at imminent risk of exposure to infection, as was done in the ring vaccination trial (Henao-Restrepo et al., 2017).

A particular challenge for vaccine trials with response-adaptive randomisation is accounting for time trends (Villar et al., 2018), in that the design needs to be robust to confounding effects from potential time trends in the data, which we would expect to be considerable for an epidemic. An important question is whether the adaptive design is beneficial given the required volume of information. Put another way, can other choices in the design space be made to increase power enough to meet the additional costs associated with adaptive designs? And is it worth it?

1.3 Addressing the challenges

Here, we use a network-based simulation to test, first, how non-adaptive design and analysis choices influence power and efficiency, and, second, whether a preferable trial design can be found through adaptation. We present two trial design and analysis elements: ring recruitment vs. random recruitment, and the use of time information for downweighting data where it’s unclear whether a participant should have been excluded from the trial. We identify choices that increase power for a fixed number of participants, in line with other discussions of optimal vaccine trial designs (Kahn et al., 2018; Nason, 2016). These results are important for vaccine trials, where participants might be sparse and the infection rate low, and for response-adaptive randomisation, where unequal allocation to trial arms comes at the cost of power. We then compare the merits of different response-adaptive randomisation designs and discuss in which circumstances the design might be favoured in terms of population benefit from the trial (Bellan et al., 2017). We focus on COVID-19, and make reference also to EVD in illustrating the design elements, pertaining to phase 2b and phase 3 clinical trials.

1.4 Outline

In Section 3 we describe our non-adaptive trial design and analysis choices. To this basis, adaptive elements are added in Section 4. We assess and evaluate the designs through simulation, with operating characteristics and summary statistics, and in Section 5 we conclude. The simulation model and trial mechanics are fully described in Appendix A, and code is available at github.com/robj411/ADAGIO/COVID19.

2 Simulation study

We use a network model to simulate an epidemic occurring in a population, in which the vaccine trial operates. The model consists of a population from which the trial participants are selected, and who are connected through a contact network, as in Hitchings et al. (2018). It is through these connections that the disease spreads, and each infected individual transitions through the stages of the disease (see Appendix A for details).

2.1 Trial design

The elements of trial design we explore are specific to an infectious disease with the dynamics and means of spread of EVD or COVID-19. In addition, our trial is for a vaccine whose time to seroconversion is fast, that is, a (hypothetical) disease-specific antibody test would go from negative to positive within 15 days, for all participants with high confidence (such as Marzi et al. (2015), <https://investors.modernatx.com/news-releases/news-release-details/moderna-announces-positive-interim-phase-1-data-its-mrna-vaccine>). We use seroconversion as a proxy for protection (a vaccine-induced protective immune response) in our simulations.

We develop a trial design sequentially by initially defining a “base case” and comparing to an equivalent design with one element changed, carrying forward the better-performing design. We make the comparisons using the model and trial system described in Appendix A. The starting “base case” for our simulations is similar to the WHO Blueprint: participants are recruited at random; cases are excluded if they display symptoms within nine days; the allocation probabilities are fixed and equal; the final follow-up time is 25 days; the trial terminates once 24 confirmed cases have been observed, and the vaccine efficacy, for an assumed leaky vaccine, is 0.7.

There are some differences between our base case and the Blueprint. We choose nine days to align with our assumed time to seroconversion and incubation period. Henao-Restrepo et al. (2017) excluded those exhibiting symptoms within 10 days in their trial for an EVD vaccine, and WHO R&D Blueprint (2020) propose 14 days for COVID-19. We recommend that the number of days is chosen on a case-by-case basis. The Blueprint uses a time-to-event outcome, whereas we use a binary outcome, in order to later enable response-adaptive randomisation using established methods. For the same reason, we have a final follow-up time, whereas the Blueprint does not.

In the comparisons we make, all designs are individually randomised and compare to a placebo control (both suggested by Kahn et al. (2020)), and use a binary outcome (which is whether or not a person has a PCR-confirmed diagnosis of COVID-19) at a final follow-up time of 25 days. We explore the effect of the recruitment schedule, the exclusion criterion (which is applied retrospectively, at the end of the trial), and the possibility of response-adaptive randomisation.

Primary endpoint We use PCR-confirmed infection as the primary endpoint, which in our simulations corresponds to any person who becomes symptomatic. PCR at the time of symptoms would be used preferentially to antibody testing due to the difficulty in distinguishing between antibodies signalling past infection and those signalling vaccination (Dean et al., 2019), and because seroconversion following infection may occur many days after onset of symptoms (Long et al., 2020). Thus we assume that through surveillance and self-reporting, symptomatic participants correspond to confirmed cases, as was the case in Henao-Restrepo et al. (2017). Some people test positive for SARS-CoV-2 but show no symptoms. In our simulation, we include these people as infectious but not symptomatic. As we employ case confirmation and weighting on the basis of symptoms, these cases are omitted from our analyses: they are counted as successes, rather than fails, in terms of the binary endpoint.

Follow-up time The single final follow-up time of 25 days after randomisation is similar to the follow-up time of 21 days used for EVD (Hitchings et al., 2017; Henao-Restrepo et al., 2015). While they chose a follow-up time of less than one month in order to then vaccinate the control arm, we make a similar choice in order to enable response-adaptive randomisation.

Our final follow-up time of 25 days is used to define the primary endpoint both for adapting (as described in Section 4.1) and for the final analysis at the end of the trial. We fix the follow-up time for reasons of simplicity and computational efficiency, in that we simulate a network only up to day 25.

How large should the trial be? We use an event-driven (Schoenfeld, 1983) (or information-based (Mehta et al., 2009)) trial design. That is, the trial is terminated following the accumulation of a prespecified number of confirmed cases, rather than a prespecified number of enrolled participants, corresponding to a fixed sample size. (N.B. when confirmed cases are weighted, the number of cases becomes the effective number of cases, and the sample size the effective sample size.) See Appendix D for a presentation of the rationale and consequence of this choice.

2.2 Operating characteristics and metrics reported

The results we report in these comparisons represent 1,000 (10,000) simulated trials (unless otherwise stated), where each “trial” involves independent networks (as many independent networks as required to achieve a particular total effective number of cases), or a fixed number of independent networks, as stated. We report operating characteristics including the number of people enrolled and the number of those symptomatic (which we equate to a confirmed case), the power, the estimated vaccine efficacy (VE) – the percent reduction in attack rate for vaccinated people compared to unvaccinated people, and the type 1 error rate, alongside the details of the design. The duration of the trial is reported in days assuming an average of 32 participants enrolled per day.

The power (the probability to correctly reject the null hypothesis) and type I error rate (the probability to incorrectly reject the null hypothesis) are estimated as the proportion of simulations under the positive effect and the null, respectively, for which the null hypothesis was rejected (see Appendix B for mathematical details). The vaccine efficacy is estimated using all simulations under the positive effect (whether or not the trial realisation concluded efficacy). Other results presented (e.g. the numbers of people enrolled and confirmed as cases) are computed under the alternative unless stated otherwise.

Additionally, we report the “prevented export infections”, defined as the expected number of infection events of people not in the index case’s contact network, for 100 contact networks with no vaccine effect, minus the equivalent number given a vaccine effect of 0.7. While we don’t expect this to be predictive of actual numbers of infections observed, the relative numbers between methods are indicative of the trial’s possible or probable effects on the wider epidemic, which would be a useful feature to develop in order to enable assessment of the impacts of trial designs.

3 Non-adaptive elements and analysis choices

3.1 Recruitment schedule: contact tracing

3.1.1 Methods

We compare random recruitment to the ring-recruitment strategy employed in Henao-Restrepo et al. (2017), which was inspired by the role of ring vaccination in international elimination of smallpox (Fenner et al., 1988). Participants are eligible for enrolment when someone in their contact network is confirmed as a “case”, starting from the first day of the trial.¹ We define the ring as consisting of contacts and contacts of contacts, where a contact-of-contact is to a contact as the contact is to the index case. These contacts are found through “contact tracing”, as described in ECDC (2020). In the terminology of Fyles et al. (2020), our rings, consisting of two levels of contacts, come from “two-step contact tracing”. In the terminology of graph theory, the “ring” is the neighbourhood of order 2 of the index case. Recruiting participants who are at imminent risk of infection is akin to the method of recruitment of high-risk participants, which was the motivation for Samai et al. (2018) to recruit front-line health workers in the Ebola outbreak of 2014–2016.

We distinguish between different types of relationships (or edges) in the network. We define a “predictable” edge as a relationship that will be discovered through contact tracing. The notion of unpredictable edges here is related to “environmental transmission” defined in Ferretti et al. (2020), that is, through the environment and not from a close contact. The notion of a “predictable” edge here is like the definition of an “acquaintance” in Kucharski et al. (2020) as someone who has been “met before”. We define the “predictable fraction” of a network as $r_s = e_s w_s / (e_s w_s + e_u w_u)$, the weighted sum of predictable edges over the total weighted sum of edges, where there are e_s predictable and e_u unpredictable edges, with weights w_s and w_u respectively. The edge weight is a property of the network that we use to convey the extent of contact, which translates to the relative propensity to transmit infection through this contact. In our work, the predictable fraction is entirely determined by how “contact tracing” is implemented: any contact that is not traced is by definition unpredictable. For our simulation, we specify three types of edge (household, workplace, and random) and stipulate that household and workplace edges are predictable and have weight 1, while random edges are unpredictable and have weight 1/10. (See Appendix A for details of the network.)

¹Note that here we use “ring” with reference to the method for recruitment. In our work this is distinct from the randomisation method, which, in other implementations of the strategy, has been cluster randomisation (Henao-Restrepo et al., 2015).

To assess the value of using ring recruitment through contact tracing for recruiting from a group at imminent risk, first, we compare the ring recruitment method to a random recruitment method (where the predictable fraction is 0.91, which is the value from our standard network; see Appendix A). We also demonstrate the relationship between the predictable fraction and the trial’s power. We adjust the predictable fraction by changing the edge weights.

3.1.2 Results

Table 1: Comparison of designs where participants are recruited following the ring strategy vs. recruited at random. The trial follows the FR design with a follow-up time of 25 days. The trial ends when an effective number of 24 cases have been observed. Standard deviations in brackets.

Recruitment	Number of participants	Number of confirmed cases	Vaccinated	Power	Type 1 error	VE estimate	Prevented export infections
Random	16866 (3353)	40	8433	0.72	0.04	0.62 (0.19)	0.72
Ring	1725 (502)	41	862	0.71	0.04	0.59 (0.19)	5.85

Table 2: The effect of the predictable fraction of transmission events on power. The trial follows the FR design with a follow-up time of 25 days and ends once 100 contact networks have been enrolled. Participants are recruited following the ring strategy. Standard deviations in brackets.

Predictable fraction	Number of participants	Number of confirmed cases	Vaccinated	Power	Type 1 error	Number of confirmed cases (VE=0)	VE estimate
1.0	2907 (168)	83	1452	0.92	0.04	106	0.59 (0.13)
0.8	2909 (172)	62	1453	0.80	0.04	77	0.58 (0.17)
0.6	2904 (168)	45	1450	0.65	0.05	54	0.57 (0.23)
0.4	2899 (178)	30	1448	0.42	0.05	36	0.5 (0.35)
0.2	2914 (169)	19	1456	0.27	0.06	21	0.4 (0.7)

In Table 1, where recruitment is random, rather than through contact tracing, we have to recruit many more participants in order to observe the requisite number of events. In addition, many more people in the general population must become confirmed cases, and very many fewer exported infection events are prevented, in our random, non-targeted trial population.

In Table 2 we consider five possible values for the “predictable-to-unpredictable contacts” balance by holding the edges fixed and adjusting the edge weights. To show power as a function of the predictable fraction, we fix the trial size to be 100 contact networks of participants. The results show the extent to which power diminishes as the fraction of predictable edges decreases. As contacts become more “predictable” (thus more contacts are recruited into the trial), the number of confirmed cases within the trial increases, especially in the “null case” where there is no vaccine effect (VE=0). This is a result of the recruitment method better targeting those at risk, as the effective vaccine will protect more of those at risk when those at risk are more concentrated within the contact network enrolled.

3.1.3 Interpretation

Where it is possible to predict who is at imminent risk of infection, ring recruitment designs are advantageous for vaccine trials. We expect to see a high predictable fraction for diseases whose transmission depends on exchange of or exposure to bodily fluids, such as HIV and EVD, as well as for COVID-19 in societies under “lockdown”, where public spaces are closed and people stay at home, and/or there is extensive quarantining. We expect that the lesser the restrictions to movement and activity, the lesser the fraction of predictable transmission events. We therefore carry forwards this design as the “base case” for further comparisons.

3.2 Exclusion criterion implemented at analysis points

In a vaccine trial, one of our aims is to evaluate vaccine efficacy. Our target population is the general population, and, using ring recruitment, our trial population are people at imminent risk of infection. Therefore, to learn the vaccine efficacy for a general person, we exclude from the analysis anyone diagnosed within nine days of randomisation, assuming they were infected either before randomisation or before the vaccine has a chance to take effect.

Although when vaccine efficacy is defined, in general, no explicit mention is made of the time to seroconversion, we assume that its consideration is implicit. Where efficacy is defined as the reduction in infection risk between the vaccinated population and the unvaccinated population (Weinberg and Szilagyi, 2010; Shim and Galvani, 2012), we would add the following clause: assuming that the vaccinated population have reached the maximum state of protection they are capable to reach. We equate this state with seroconversion. We assume that this is implicit in, for example, studies that conclude efficacy for vaccines whose instructions for correct usage includes a time component, i.e. any vaccine that is advertised as having a particular efficacy and requiring a certain number of days to take effect (following the same process as, for example, meningococcal meningitis vaccination, which should be sought two to three weeks before travel (<https://www.nhs.uk/conditions/travel-vaccinations/jabs/>)).

By excluding observations, we lose some information. Therefore we seek other means to target our ideal endpoint, which is: whether or not a participant became infected after randomisation and vaccine-induced seroconversion, where participants who became infected before are excluded from the trial. For diseases such as EVD and COVID-19, we cannot know the day on which a person is infected, and we cannot know the day on which a vaccinated person seroconverts. However, we can observe whether or not a person is symptomatic, at which point infection can be confirmed through PCR. Therefore, when we undertake the analysis at the end of the trial, we supplement the primary endpoint – the observed infection status at day 25 – with a retrospective exclusion criterion: we use our knowledge of the disease (e.g. the incubation period) and the day a person becomes symptomatic to weight their inclusion in the analysis. If the incubation time is the estimated time to fever, for example, then the participant (or a person close to the participant) should be asked on which day the fever started. In addition to these factors, which were considered in Henao-Restrepo et al. (2015), one could include also a “reporting time” if it is expected that some time elapses between symptom onset and the reporting of symptoms.

In summary, as we cannot know when someone was infected, we propose a primary endpoint for the final analysis defined as whether or not a person is symptomatic and, in addition, the day of symptom onset relative to randomisation, which informs a retrospective exclusion criterion. This exclusion of participants manifests as a “weight” between 0 and 1, which we compute as part of analysing the results so far, and apply “retrospectively”, as if we had excluded a participant from the beginning. This challenge is specific to vaccine trials with recruitment of people at imminent risk; the primary measure is usually observed with confidence when participants enter into a treatment trial, and, in vaccine trials with participants recruited at random, the likelihood for a participant to develop symptoms soon after enrolment is small, and thus the information cost of binary exclusion is small.

3.2.1 Methods

The probability that a person whose symptoms began after vaccination (and corresponding seroconversion) was infected after seroconversion depends on the effect of the vaccine. If the vaccine is effective, then this probability is smaller than it would be if the vaccine had no effect. Therefore, as an additional development, we estimate the vaccine efficacy and the weights together iteratively, so that both the retrospective exclusion weights and the vaccine efficacy estimate are updated as results accumulate (if, as in a response-adaptive design, we were to calculate them as results accumulate). See Figure 1 for illustration, and Appendix B for the derivation. We assess this weighting method in terms of power and type 1 error compared to binary exclusion. See Appendix B for their mathematical presentation.

3.2.2 Results

In Table 3, we compare the two methods for defining the retrospective exclusion criterion. By downweighting (“Continuous”) inclusion, rather than applying a binary rule, there is an increase in power, and the VE estimate is closer to the true value of 0.7. The gain in power is in large part due to the correct accounting for the vaccine efficacy when determining which early cases are likely to have been infected before randomisation. The other operating characteristics are similar between the methods.

3.2.3 Interpretation

The method we present for the final analysis achieves a higher power and more accurate estimate of the vaccine efficacy for a similar trial size. It is specific to the design of vaccine trial we consider, where it is not clear at randomisation who is eligible for the trial, nor if the vaccination will result in immunity before a vaccinated




Figure 1: The curves show the probability that the value on the x axis was the day a person became infected, given the day they showed symptoms, whether or not they were vaccinated, and the effect of the vaccine (where we omit seroconversion for simplicity of presentation). The inclusion weight given the day of symptom onset is the area under curve to the left of the randomisation day. Two participants show symptoms on the same day. One (turquoise) is vaccinated. The other (navy) isn't. Given the distribution of the incubation period, the navy person has a probability of 0.83 of having been infected after randomisation day. Their weight is therefore 0.83. Given the vaccine efficacy estimate at this particular point exceeds zero, the turquoise person is less than 83% likely to have been infected after randomisation day. In the binary-weighting case, both these participants have weight 0 (if day of symptom onset is less than the threshold of nine days) or weight 1 (otherwise).

Table 3: Comparison of designs where the binary endpoint has a binary weight or a continuous weight. Participants are recruited following the ring strategy. The trial follows the FR design with a follow-up time of 25 days. The trial ends when an effective number of 24 cases have been observed for the continuous weight and 26 for the binary weight, in order to achieve comparable trial sizes in terms of the number of participants. Standard deviations in brackets.

Weighting	Number of participants	Number of confirmed cases	Vaccinated	Power	Type 1 error	VE estimate	Number of participants (null)	Prevented export infections
Binary	1866 (533)	44	933	0.74	0.04	0.59 (0.18)	1176 (383)	5.62
Continuous	1917 (541)	45	958	0.81	0.04	0.65 (0.18)	1158 (383)	5.41

participant is exposed to infection. Our analysis uses information pertaining to disease and vaccination timings in order to weight confirmed cases in the analysis.

3.3 Conclusions

As vaccine trials risk being underpowered, due to rarity of events and due to difficulties in recruiting participants, and as response-adaptive design choices targeting patient benefit come at the cost of power, we explored and presented design and analysis choices that can be used to mitigate these effects and bolster power. The design choice is to recruit contacts at imminent risk of infection. The analysis choice is the use of a continuous value to weight cases observed for inclusion in the efficacy calculation and significance test. This analysis choice goes hand in hand with the design choice of ring recruitment, which unavoidably will enrol a number of participants who are already exposed to the disease. By condensing the gathering of information to a short time horizon, this design enables the exploration of response-adaptive randomisation in the context of a vaccine trial, to which we turn now.

4 Response-adaptive randomisation

In the response-adaptive trial simulations, at fixed-frequency, pre-specified moments in the course of the trial, we update the allocation probabilities, using the formulae presented in Sections 4.1.1 and 4.1.2. The allocation probabilities are the probabilities of individuals being allocated to each arm of the trial. We choose the frequency of adaptations to coincide with the final follow-up time – that is, 25 days – and all data accrued up to the adaptation day (up to the final follow-up time of 25 days post randomisation) are used in computing the probabilities. Therefore, the randomisation ratio is not changed after each patient but rather after groups of patients are vaccinated and followed up.

4.1 Methods

To compute the allocation probabilities, first we define π_v as the probability of being allocated to arm v , where v is 0 for control and 1 for experimental. Below we present methods to compute π_1 from two broad families: frequentist and Bayesian. The methods all require first generating running estimates \hat{p}_0 and \hat{p}_1 of p_0 and p_1 , the probability of being uninfected up to the final follow-up time for the control and experimental arms, respectively. We compare all methods to a design with fixed and equal allocation probabilities, denoted FR (fixed randomisation), in which $\pi_0 = \pi_1 = 0.5$ for all time.

4.1.1 Frequentist response-adaptive

In the frequentist framework, the probabilities are calculated as the maximum-likelihood estimate (MLE): the probability is the number of successes over the total number of observations; $\hat{p}_v = 1 - f_v/N_v$, where f_v is the effective number of confirmed cases for arm v , and N_v the effective sample size for the arm (Hu and Rosenberger, 2006).

In the Rosenberger et al. method (Rosenberger et al., 2001), favourable outcomes are optimised subject to power. To use this method, we define $\rho = \sqrt{\frac{\hat{p}_1}{\hat{p}_0}}$ as the allocation ratio of experimental to control. Then the allocation probability to the experimental arm is

$$\pi_1 = \frac{\rho}{\rho + 1} = \frac{\sqrt{\hat{p}_1}/\sqrt{\hat{p}_0}}{\sqrt{\hat{p}_1}/\sqrt{\hat{p}_0} + 1} = \frac{\sqrt{\hat{p}_1}}{\sqrt{\hat{p}_1} + \sqrt{\hat{p}_0}}.$$

For example, if the infection rate for the control arm is 80% ($\hat{p}_0 = 0.2$) but a vaccine can reduce that to 20% ($\hat{p}_1 = 0.8$) then the Rosenberger et al. method would allocate patients in a 2:1 ratio favouring the vaccine ($\rho = \sqrt{4} = 2$, and $\pi_1 = 2/3$). This method is designed to balance power and patient benefit, in terms of maximising treatment successes. In the multi-arm case, worse-performing experimental arms can be deprioritised, enabling prioritisation of other arms. For the two-arm case, then, we expect to see little overall gain, as there is only one experimental arm.

In the Neyman method we use the Neyman frequentist rule, which is designed to improve power compared to equal allocation (Rosenberger et al., 2001), by setting

$$\pi_1 = \frac{\sqrt{\hat{p}_1(1 - \hat{p}_1)}}{\sqrt{\hat{p}_0(1 - \hat{p}_0)} + \sqrt{\hat{p}_1(1 - \hat{p}_1)}}.$$

For $\hat{p}_0 = 0.2$ and $\hat{p}_1 = 0.8$, this rule has $\pi_1 = 0.5$. (For the case $\hat{p}_0(1 - \hat{p}_0) = 0$ or $\hat{p}_1(1 - \hat{p}_1) = 0$, we set $\pi_0 = \pi_1 = 0.5$.)

4.1.2 Bayesian response-adaptive (Thompson sampling)

At the adaptation point, we obtain the posterior distribution of p_i given a uniform prior and the observed data as $\text{Beta}(1 + N_v - f_v, 1 + f_v)$. The allocation probability π_i is defined in terms of these posterior probabilities, computed by sampling, as (Thall and Wathen, 2007)

$$\pi_1 = \frac{\Pr(p_1 > p_0)^\phi}{\Pr(p_1 > p_0)^\phi + \Pr(p_1 < p_0)^\phi},$$

where we define a tuning parameter ϕ . For Thompson sampling (TS), we set $\phi = 1$, so π_1 is just $\Pr(p_1 > p_0)$, and for TS with tuning (TST), $\phi = j/e$, with j the day of the current update, and e the trial's expected total duration, e.g. 100 days. One might instead choose to adapt according to number of cases seen, so that $e = 24$ effective cases. ϕ therefore takes the value 0 at the beginning of the trial and goes to 1 as the trial progresses, so that the speed with which the allocation probability can reach extreme values (0 or 1) is tempered.

The Thompson sampling methods have a possibility of generating very high allocation probabilities. While tuning slows this process, we find that the TST method tends to 1 over a few adaptations (results not shown).

We therefore bound the allocation probability above at 0.8 and below at 0.2 for all methods. We additionally (but separately) consider terminating the trial and concluding efficacy if we compute an allocation probability of 0.99 (Brueckner et al., 2018).

Time trends The epidemic unfolding in real time gives rise to time trends in incidence of the disease, which in turn will result in patient drift in the trial (Proshan and Evans, 2020; ?). By “patient drift”, we mean some effect that changes over time and influences all participants regardless of vaccination status (as opposed to “temporal trends”, which we define to be a time-dependent effect that will affect some but not all arms in the trial, and which we do not consider here). For instance, a natural increase or decline in incidence, or a step change due to government policy on social contact, as proposed by Ferguson et al. (2020), or even a change in the recruitment strategy, would induce a time trend in the trial data. As both the adaptive trial design and the epidemic are changing over time, in analysing the results, we account for time dependencies of disease exposure in order to infer the effect of the experimental vaccine.

We use randomisation-based inference to correct for patient drift as described by Simon and Simon (2011), which involves resampling the data in order to generate a new null distribution for the test statistic to which to compare the one we compute. We present the resulting powers and rates of type 1 error. In the Appendix, we demonstrate that the correction preserves type 1 error at the cost of a little power in the context of a strong downward linear trend in incidence. Other methods include a simple stratified analysis of the data (Karrison et al., 2003; Chandereng and Chappell, 2019).

Table 4: Comparison of response-adaptive designs. The outcome has a continuous VE-corrected weighting. Participants are recruited following the ring strategy. The final follow-up time is 25 days. The trial ends when 24 effective cases have been observed. Standard deviations in brackets. Correction for time trend uses the resampling method of Simon and Simon (2011).

Adaptation	Number of participants	Duration (days)	Number of confirmed cases	Vaccinated	Power	Power (corrected)	Type 1 error	Type 1 error (corrected)	VE estimate	Prevented export infections
Ney.	1757 (463)	79 (14)	49	749	0.81	0.74	0.06	0.05	0.66 (0.2)	4.64
Ros. et al.	1916 (534)	84 (17)	51	965	0.80	0.77	0.05	0.04	0.64 (0.19)	5.18
TST	2238 (712)	94 (22)	56	1423	0.76	0.70	0.04	0.04	0.62 (0.19)	6.33
TS	2346 (745)	98 (23)	57	1582	0.74	0.47	0.04	0.06	0.62 (0.19)	6.60
FR	1906 (524)	84 (16)	51	953	0.80		0.05		0.64 (0.19)	5.28

Table 5: Thompson-sampling response-adaptive designs where the trial is terminated when the allocation probability reaches 0.99. The outcome has a continuous VE-corrected weighting. Participants are recruited following the ring strategy. The final follow-up time is 25 days. The trial ends when 24 effective cases have been observed. Standard deviations in brackets. Correction for time trend uses the resampling method of Simon and Simon (2011).

Adaptation	Number of participants	Duration (days)	Number of confirmed cases	Vaccinated	Power	Power (corrected)	Type 1 error	Type 1 error (corrected)	VE estimate	Prevented export infections
TST	2032 (638)	88 (20)	51	1261	0.77	0.76	0.04	0.04	0.64 (0.19)	6.62
TS	1799 (740)	81 (23)	45	1148	0.80	0.74	0.04	0.05	0.67 (0.21)	5.99

11

Table 6: Comparison of response-adaptive trials that last at most 85 days. We compare their profiles in terms of the totals vaccinated up to day 100. “Vaccinated in trial” is the expected number vaccinated during the trial, and “Vaccinated up to day 100: assuming 32 (320) per day” is the expected number vaccinated after the trial up to day 100, estimated as 32 (320) people vaccinated per day, for the days remaining after the end of the trial, multiplied by the power (the probability to have concluded efficacy and rolled out the vaccine). Standard deviations in brackets. Correction for time trend uses the resampling method of Simon and Simon (2011).

Adaptation	Sample size	Duration	Number of confirmed cases	Vaccinated in trial	Power	Power (corrected)	Export infections in trial	Vaccinated up to day 100: assuming 32 per day assuming 320 per day	
Ney.	1939 (115)	85 (4)	55	807	0.81	0.76	25	1177	4474
Ros. et al.	1939 (115)	85 (4)	52	977	0.81	0.79	24	1360	4774
TST	1939 (115)	85 (4)	48	1232	0.74	0.69	23	1567	4558
TS	1939 (115)	85 (4)	47	1264	0.70	0.54	23	1527	3875
FR	1939 (115)	85 (4)	52	970	0.81	0.81	24	1363	4868

4.2 Results

The results for a suite of adaptive trials are shown in Table 4. We report the power corrected for the time trend in the epidemic following Simon and Simon (2011). See Figure 5 for operating characteristics under different trends and Figure 4 for trajectories of allocation rates.

In Table 4, we see that the Thompson sampling methods (TST and TS) allocate more participants to vaccination instead of control. As a result, there are fewer infections exported from the network, and the power is lower. The Neyman method has lower power also because of unequal allocations, in this case with more participants in the control arm, which incurs a cost in power to correct for patient drift. The Rosenberger et al. method is most similar to the fixed and equal randomisation design, matching it in terms of number of participants, participant allocation and power, and vaccinating slightly more participants.

In Table 5, we consider the Thompson sampling methods again, where we include the possibility that the trial concludes efficacy and terminates should the calculated allocation probability reach 0.99. Doing so results in much shorter trials: on average 6 and 17 days shorter (for TST and TS, respectively) than the trials that do not stop early, and, additionally, the TS trial is shorter on average than the fixed-randomisation trial. We note that these designs do not suffer as much the penalty to correct for patient drift as the trial does not spend such an extended period with unequal allocations.

The trade-off between power, number of people infected, and the time taken to conclude efficacy is the crux of this design choice. In Table 4, we fixed the number of cases observed in the trial population to a total weight of 24, which sets the health cost to the trial participants. We can then trade off the number of prevented exported infections against time to conclude directly: the fastest trial prevents fewest export infections; the longest trials prevent most export infections. However, the trials with the highest power are those with intermediate duration and intermediate prevented export infections.

Another means to compare the methods would be to ask what can the trial achieve within, say, 100 days, assuming an efficacious vaccine, and where the trial lasts at most 85 days. We assume that the trial begins when the first participant is enrolled and ends when last participant has been followed up for 25 days. Assuming one contact network can be enrolled per day, 60 contact networks are enrolled. Table 6 shows this comparison and that, when set up with a limited time, the trials that allocate most participants to the efficacious experimental arm – TS and TST – would be preferable in terms of successfully vaccinating the most people during the trial, and up to day 100, if the rate of vaccination after the trial is the same as the rate of enrolment into the trial. (There are also fewer export infection events in the course of the trial – both in total and per day.) However, if, after the trial, we can vaccinate ten times as many people per day as were enrolled in the trial, the designs with higher powers look better, as they are better able to identify an efficacious vaccine, and therefore will vaccinate more people in the period immediately after the trial.

Together, Tables 4 and 6 suggest that where participants are abundant, a trial can justify, in terms of power, allocating more participants to the best performing arm, which would imply vaccinating more people (if the vaccine seems promising from data so far). Where participants are abundant, the epidemic is spreading, and so vaccinating more people would have the desirable side effect of attenuating the epidemic to some extent. On the other hand, where the trial is limited by the number of participants it can recruit, as it would be in a declining epidemic, a trial that maximises power might be preferable. Such a design might not prioritise vaccination, based on what we see in Tables 4 and 6. The rationale is that, even if the trial makes fewer vaccinations within the trial, it will result in more vaccinations overall, as an efficacious vaccine is more likely to be identified by prioritising information gain. This would be relevant for a trial that operates in locations where the epidemic has subsided considerably but there exist other locations that still have high incidence. In this scenario, there is little to gain by prioritising vaccination over information.

5 Conclusion

We have considered designs for vaccine trials in epidemic outbreaks, focusing on highly infectious, person-to-person-spread disease, using the example of COVID-19. Because the disease spreads through person-person contacts, it is possible to anticipate who is at imminent risk of becoming infected. Therefore we use the ring recruitment strategy for this disease.

We presented a trial design that works well if the time to trace contacts and the seroconversion time together have duration of order no more than that of the spread of the disease. Otherwise, the epidemic would outrun contact tracing and therefore the trial. Thus trials for vaccines that require two shots are precluded, as are trials for diseases that spread faster than contacts can be traced, for example.

Our simulation differs from the ring vaccination trial for EVD as described in Henao-Restrepo et al. (2015) in a number of ways. We simulate an individually randomised controlled trial (iRCT), whereas their trial was cluster randomised (cRCT). One consequence of this is they had stricter criteria for eligibility: in a cRCT, a person is enrolled only if their cluster is enrolled. In our simulation, each person’s enrolment is independent. So, a person identified in a contact network with 60% or more overlap with an enrolled contact network would not be enrolled in a cRCT, but they can be enrolled in an iRCT. In general, we expect the methods we presented to be compatible with a cRCT design.

Both Henao-Restrepo et al. (2015) and WHO R&D Blueprint (2020) use a time-to-event endpoint, whereas we use a binary endpoint. The Blueprint employs random recruitment and individual randomisation, to which we compared ring recruitment with individual randomisation.

We have discounted asymptomatic people, assuming that case confirmation follows symptoms. These participants could be confirmed as “cases” if trial participants are routinely PCR tested. Regular PCR testing of all participants would boost power for all trial designs, but might not be operationally feasible. We do not anticipate that their omission will bias the results as they are randomly allocated to arms in the same way as other participants. Their inclusion would, however, change the endpoint from “disease” to “infection” (Hudgens et al., 2004). These endpoints might differ only in terms of the power given a trial size; they might also differ in vaccine efficacy, depending on how the vaccine operates.

5.1 Design and analysis choices

We compared random recruitment to ring recruitment and identified a means to anticipate the utility of recruiting those at imminent risk of infection. We demonstrated that, in agreement with intuition, ring recruitment will be advantageous when there is a contact-tracing protocol available that will identify those at imminent risk. A consequence of recruiting people at imminent risk of infection is that response-adaptive designs are enabled, which require that the outcome is observed soon after randomisation relative to the duration of the trial.

The utility of the ring design for recruitment for vaccine trials for infectious diseases depends most on the proportion of contacts of an infectious person that can be traced (see Table 2). The higher the predictable fraction, the more cases there will be per participant, and the higher the power. In order to power the trial in the first instance (and perhaps to determine if the ring is an appropriate recruitment strategy), an estimate must be made of the balance of predictable to unpredictable transmission events. This might be informed by consideration of data from early efforts of contact tracing: one might ask what proportion of events might have been anticipated, given relationships to prior events, and what sorts of relationships were important. This can be further informed by changes in social structure: e.g. if many predictable events occurred in particular locations, which since have been closed, we can estimate the resulting effect on the predictable-to-unpredictable balance.

Estimation of the predictable fraction could be informed by consideration of available data, e.g. for our example, where we consider household and workplace contacts, we would estimate the predictable fraction as the proportion of new cases that are linked through household and workplace contacts to known, pre-existing cases, or transmission events could be identified following tracing of all contacts (COVID-19 National Emergency Response Center and Epidemiology Case Management Team Korea Centers for Disease Control Prevention, 2020). Alternatively, or additionally, we could use data on contact patterns (Klepac et al., 2020). If we consider that the power associated with our estimated predictable fraction is too low, we could consider adding additional contacts to the network, for example neighbours. If it is beyond the operational capacity of the trial to trace sufficient contacts to achieve a certain desired power, then the ring recruitment strategy is unlikely to deliver an adequately powered trial.

We presented a new method to analyse the data at the end of the trial, applying continuous weights to the data according to the number of days after randomisation on which symptoms first appear. The approach is similar to that used by Henao-Restrepo et al. (2017), whose exclusion criterion corresponds to a binary weighting.

We use instead a continuous weighting and, in addition, we estimate the vaccine efficacy simultaneously. The effect is an increase in power per participant.

If, for the exclusion criterion, the choice was made that all cases were included (i.e., there is no exclusion, and no temporal data are collected), we would be effectively testing “what is the effect of the intervention”, which is less than the effect of the vaccine, as some people vaccinated cannot benefit from the vaccine. This quantity might be of interest in its own right, for example to test what the effect might be of vaccinating in addition to contact tracing as a method for containment of the outbreak or an “intention to treat” design (Fyles et al., 2020), establishing the “vaccine effectiveness” (as opposed to vaccine efficacy) in those conditions (<https://www.cdc.gov/flu/vaccines-work/effectivenessqa.htm>). However, for the quantity we are interested in – the vaccine efficacy – this measurement is not a good proxy.

5.2 Adaptations

We explored four different response-adaptive designs in the context of the vaccine trial. We found that there is a penalty in power of the designs due to the requirement to account for patient drift. The more the allocation deviates from equality, the greater the penalty. Bounding the allocation probabilities between 0.2 and 0.8 protects the designs from very severe penalties, as well as making the designs more realistically acceptable to triallists.

We propose response-adaptive randomisation as a means to achieve some population-health objective, which could be attenuating infection spread, aiming to conclude the trial quickly and therefore release an efficacious product for use, or a combination of the two, subject to a particular power. Alternatively, one might consider the balance between power and number of people vaccinated, subject to a fixed trial duration.

For designs that terminate the trial after a prespecified number of cases have been observed (controlling the power), the “cost” to the trial participants, measured in terms of number of cases (and therefore “treatment failures”) is fixed, so the health impact of the design choice can be measured by the impact on the rest of the population. This in turn will depend on the local current circumstances of the epidemic. Informally, where the epidemic is growing, maximising vaccination coverage will take priority, as the wider impact of the trial will be to stem the rate of transmission to the outside of the trial population. The cost to power can be met as there will likely be enough participants. Where the epidemic is attenuating, containment is not a primary concern of the trial, but attenuation of trial participants is a risk, so the more ethical choice might be the one that observes the requisite number of cases as fast as possible, in order to license the vaccine for use in other settings. Future work, including modelling a whole population-level epidemic, should quantify this formally, as in Bellan et al. (2017).

This cost to power arises from the application of a method to correct for potential patient drift. We used the resampling method of Simon and Simon (2011). We note that the cost is significant and increases as inequality in allocation increases. We suggest then that one way to improve the prospects of response-adaptive randomisation for vaccine trials would be to develop methods for correcting for the temporal trend of the epidemic, for example by stratification (Chandereng and Chappell, 2019), or by modelling incidence directly.

We recognise that the two-arm case, with one control arm and one experimental arm, gives limited insight into what a response-adaptive design can offer. In general, in a two-arm trial, one can only increase power (for a fixed number of participants) at the expense of patient benefit (within the trial) (Williamson et al., 2017; Villar et al., 2015). We note that even in this case the designs demonstrate a range in behaviours across the spectrum of power vs. patient benefit, even if none exceeds the fixed-randomisation design in this respect. We therefore look forward to multi-arm trials that show the benefits that can be brought in both vaccination successes and efficiency, motivated by the demonstration that the two-arm examples can trade one for the other in different ways.

We note that the type 1 error is not much inflated for the Thompson sampling methods, which we would expect ordinarily. This is because, under the null (i.e. no vaccine effect), little of the trial operates under an unequal allocation regime. As is shown in Figure 4, under the null, the trial typically observes 24 effective cases soon after the first adaptation day, so that the results are not very skewed. For a design that adapts earlier relative to its end time, we would expect to see an inflated type 1 error.

5.3 Follow-up time

It can be a challenge with these trial designs to choose appropriate follow-up times. These will be informed by some extent by the parameters governing disease transmission dynamics, and also on the ability to trace contacts, the effective reproduction number, and other quantities that are difficult to know at the outset. Extensive

simulation to test a range of values is necessary to narrow down the range of follow-up times that would suit the design objectives.

We use a single final follow-up time, 25 days. This was chosen to be large enough to observe cases with confidence that the results were meaningful, in that the case became infected after enrolment and after seroconversion, and small enough so that accumulated data could be used to adapt the trial in a way that delivered on some other objective, e.g. to maximise efficiency or vaccination successes.

Consistent with all of the above choices would be (a) to follow up all participants beyond the stated number of days: that is, each participant has one endpoint for adaptation and one endpoint final analysis, which could be the outcome at end of the trial. This will increase power for all designs, as more information would accumulate over time, though Hitchings et al. (2017) show that, in a declining epidemic, where the effective reproduction number is below 1, there is a limit to the return on extending the follow-up time for a ring-recruitment vaccine trial. The other option, (b), would be to vaccinate the control arm at the final follow-up time. Then the statistical comparison between the two groups terminates at this point (so there is one endpoint only) but there is an ethical benefit of having potentially vaccinated more people. This was the approach used in two EVD trials (Henao-Restrepo et al., 2017; Samai et al., 2018).

5.4 Endpoint

Throughout our analyses we have used a binary endpoint weighted by a retrospective exclusion criterion. A time-to-event-type analysis (TTE) would be enabled by the data we assume is collected. We use a weighted binary outcome so that the temporal element of our outcome is associated exclusively with the exclusion criterion. By weighting a binary outcome we can account more naturally for measurement error in the data, such as time from infection to symptom onset, as well as exclusion, than would be possible in a TTE analysis.

Additionally, we note that a central assumption of a TTE analysis is that of proportional hazards: that the hazard function for any two individuals is proportional at any moment in time and the implied ratio does not vary with time. With a network-based trial, we can verify that this assumption is violated almost certainly: a person’s hazard depends on the infection status of their neighbours. The hazard ratio between two people would then invert if one person has an infectious contact soon after their enrolment but not at the end of their follow-up time (e.g. if they are a contact of the index case), and the other person had no infectious contacts soon after enrolment but did at the end of their follow-up time (e.g. if they are a contact of a contact, and their contact becomes symptomatic during the trial).

We suggest that TTE analysis would be enabled for a trial such as this one through parametrisation of the “time” by exposure data, quantified as the cumulative time spent with a contact (or contacts) in the infectious state, rather than calendar time from the day of randomisation. This might be more easily quantified for a disease such as EVD, where symptoms correspond to risk of infection and so knowledge of the duration of symptoms of a person’s contacts enables quantification of their accumulated risk. Phone apps developed for COVID-19 could provide similar data towards parametrising such a model.

5.5 Vaccine efficacy estimates

In our simulations we consistently underestimate the vaccine effect, which were maximum-likelihood estimates. There is an apparent trend that the higher the power, the closer the VE estimate to the true effect of 0.7, which we would expect as higher power implies more trials in which efficacy is concluded, and where efficacy is concluded the VE estimate is likely to be higher. The response-adaptive randomisation methods have the same accuracy as the non-adaptive methods. We accounted for the the recruitment design in our estimates, which allows for people vaccinated to have been infected before vaccination, or before seroconversion. Thus we are making an isolated estimate of the vaccine efficacy after vaccine-induced seroconversion.

Our simulation model could be used also to investigate the impact of the vaccination strategy, in addition to the effect of the vaccine. Such a comparison would use a different participant weighting to isolate the effect of interest, for example having all weights equal to 1 would indicate the effect of the overall intervention on the whole eligible population. Similarly, using the same model with cluster randomisation would permit investigation of effects of herd immunity, again on the whole eligible population.

5.6 Future work

5.6.1 Relation to the epidemic

We have simulated a trial to operate within an epidemic. However, we have not fully simulated that epidemic itself, and so we have not included in our analyses quantification of the immediate impact of the trial on the wider population beyond superficial consideration of “exported infections”. So, while we have compared measures of patient benefit, these have been limited to those in the trial, rather than those outside. We considered the prevented export infections as a means to compare different designs, which might form a starting point for assessing the wider impact of the trial on the epidemic. However, we defer answering that question directly to later work.

Explicit inclusion of the trial in an epidemic would also permit a realistic assessment of the impact of heterogeneity in disease state of participants, of contact networks that interact or intersect, and realistic time trends. Then the likelihood of running out of eligible trial participants could be assessed and included among criteria for trial design assessment, and different paces of trials could be considered: we supposed that one contact network was initiated per day. In an epidemic, there might be the possibility to recruit multiple contact networks per day, which would impact on the characteristics of the adaptive designs.

5.6.2 Formal comparison between response-adaptive randomisation and early stopping

Having presented methods for response-adaptive randomisation in vaccine trials, the next step is to make a formal comparison between response-adaptive randomisation and strategies of early stopping, such as those presented in Henao-Restrepo et al. (2015); WHO R&D Blueprint (2020), and multi-arm, multi-stage designs. It would be pertinent to assess also what gains can be made by combining the two.

References

- Bellan, S. E., Pulliam, J. R. C., Van Der Graaf, R., Fox, S. J., Dushoff, J., and Meyers, L. A. (2017). Quantifying ethical tradeoffs for vaccine efficacy trials during severe epidemics. *bioRxiv*.
- Box, G. E. P. (1954). Some theorems on quadratic forms applied in the study of analysis of variance problems, I. Effect of inequality of variance in the one-way classification. *The Annals of Mathematical Statistics*, 25(2):290–302.
- Brueckner, M., Titman, A., Jaki, T., Rojek, A., and Horby, P. (2018). Performance of different clinical trial designs to evaluate treatments during an epidemic. *PLoS ONE*, 13(9):e0203387.
- Camacho, A., Eggo, R. M., Funk, S., Watson, C. H., Kucharski, A. J., and Edmunds, W. J. (2015). Estimating the probability of demonstrating vaccine efficacy in the declining Ebola epidemic: A Bayesian modelling approach. *BMJ Open*, 5(12):1–6.
- Chandereng, T. and Chappell, R. (2019). Robust blocked response-adaptive randomization designs. *arXiv*.
- COVID-19 National Emergency Response Center and Epidemiology Case Management Team Korea Centers for Disease Control Prevention (2020). Contact transmission of COVID-19 in South Korea: Novel investigation techniques for tracing contacts. *Osong Public Health and Research Perspectives*, 11(1):60–63.
- Danon, L., Brooks-Pollock, E., Bailey, M., and Keeling, M. J. (2020). A spatial model of CoVID-19 transmission in England and Wales: early spread and peak timing. *medRxiv*, page 2020.02.12.20022566.
- Dean, N. E., Gsell, P.-S., Brookmeyer, R., De Gruttola, V., Donnelly, C. A., Halloran, M. E., Jasseh, M., Nason, M., Riveros, X., Watson, C. H., Henao-Restrepo, A. M., and Longini, I. M. (2019). Design of vaccine efficacy trials during public health emergencies. *Science Translational Medicine*, 11(499):eaat0360.
- ECDC (2020). Resource estimation for contact tracing, quarantine and monitoring activities for COVID-19 cases in the EU/EEA. Technical Report March, European Centre For Disease Prevention And Control.
- Fenner, F., Henderson, D. A., Arita, I., Jezek, Z., and Ladnyi, I. D. (1988). *Smallpox and its eradication*, volume 6.
- Ferguson, N. M., Laydon, D., Nedjati-Gilani, G., Imai, N., Ainslie, K., Baguelin, M., Bhatia, S., Boonyasiri, A., Cucunubá, Z., Cuomo-Dannenburg, G., Dighe, A., Dorigatti, I., Fu, H., Gaythorpe, K., Green, W., Hamlet, A., Hinsley, W., Okell, L. C., van Elsland, S., Thompson, H., Verity, R., Volz, E., Wang, H., Wang, Y., Walker, P. G., Walters, C., Winskill, P., Whittaker, C., Donnelly, C. A., Riley, S., and Ghani, A. C. (2020). Impact of non-pharmaceutical interventions (NPIs) to reduce COVID- 19 mortality and healthcare demand. Technical Report March, Imperial College COVID-19 Response Team.
- Ferretti, L., Wymant, C., Kendall, M., Zhao, L., Nurtay, A., Abeler-Dörner, L., Parker, M., Bonsall, D., and Fraser, C. (2020). Quantifying SARS-CoV-2 transmission suggests epidemic control with digital contact tracing. *Science*, 6936(March):eabb6936.
- Friede, T. and Kieser, M. (2002). On the inappropriateness of an EM algorithm based procedure for blinded sample size re-estimation. *Statistics in Medicine*, 21(2):165–176.
- Fyles, M., Fearon, E., and Working, M. C. (2020). Household structured contact tracing: Branching process model. Technical report.
- Gould, A. L. and Shih, W. J. (1992). Sample size re-estimation without unblinding for normally distributed outcomes with unknown variance. *Communications in Statistics - Theory and Methods*, 21(10):2833–2853.
- He, X., Lau, E. H., Wu, P., Deng, X., Wang, J., Hao, X., Lau, Y. C., Wong, J. Y., Guan, Y., Tan, X., Mo, X., Chen, Y., Liao, B., Chen, W., Hu, F., Zhang, Q., Zhong, M., Wu, Y., Zhao, L., Zhang, F., Cowling, B. J., Li, F., and Leung, G. M. (2020). Temporal dynamics in viral shedding and transmissibility of COVID-19. *Nature Medicine*, 26(May).
- Henao-Restrepo, A. M., Camacho, A., Longini, I. M., Watson, C. H., Edmunds, W. J., Egger, M., Carroll, M. W., Dean, N. E., Diatta, I., Doumbia, M., Draguez, B., Duraffour, S., Enwere, G., Grais, R., Gunther, S., Gsell, P. S., Hossmann, S., Wattle, S. V., Kondé, M. K., Kéïta, S., Kone, S., Kuisma, E., Levine, M. M., Mandal, S., Mauget, T., Norheim, G., Riveros, X., Soumah, A., Trelle, S., Vicari, A. S., Røttingen, J. A., and Kieny, M. P. (2017). Efficacy and effectiveness of an rVSV-vectored vaccine in preventing Ebola virus disease: final results from the Guinea ring vaccination, open-label, cluster-randomised trial (Ebola Ça Suffit!). *The Lancet*, 389(10068):505–518.

- Henao-Restrepo, A. M., Longini, I. M., Egger, M., Dean, N. E., Edmunds, W. J., Camacho, A., Carroll, M. W., Doumbia, M., Draguez, B., Duraffour, S., Enwere, G., Grais, R., Gunther, S., Hossmann, S., Kondé, M. K., Kone, S., Kuisma, E., Levine, M. M., Mandal, S., Norheim, G., Riveros, X., Soumah, A., Trelle, S., Vicari, A. S., Watson, C. H., Kéïta, S., Kieny, M. P., and Røttingen, J. A. (2015). Efficacy and effectiveness of an rVSV-vectored vaccine expressing Ebola surface glycoprotein: interim results from the Guinea ring vaccination cluster-randomised trial. *The Lancet*, 386(9996):857–866.
- Hitchings, M. D., Lipsitch, M., Wang, R., and Bellan, S. E. (2018). Competing effects of indirect protection and clustering on the power of cluster-randomized controlled vaccine trials. *American Journal of Epidemiology*, 187(8):1763–1771.
- Hitchings, M. D. T., Grais, R. F., and Lipsitch, M. (2017). Using simulation to aid trial design: Ring-vaccination trials. *PLoS Neglected Tropical Diseases*, 11(3):1–12.
- Hu, F. and Rosenberger, W. F. (2006). *The Theory of Response-Adaptive Randomization in Clinical Trials*. John Wiley Sons, Ltd., New Jersey.
- Huang, L., Bai, J., Yu, H., and Chen, F. (2018). Sample size re-estimation without un-blinding for time-to-event outcomes in oncology clinical trials. *Journal of Biomedical Research*, 32(1):23–29.
- Hudgens, M. G., Gilbert, P. B., and Self, S. G. (2004). Endpoints in vaccine trials. *Statistical Methods in Medical Research*, 13(2):89–114.
- Kahn, R., Rid, A., Smith, P. G., Eyal, N., and Lipsitch, M. (2018). Choices in vaccine trial design in epidemics of emerging infections. *PLoS Medicine*, 15(8):1–12.
- Kahn, R., Villar, S. S., and Lipsitch, M. (2020). Innovative Vaccine Trial Designs for EID Outbreak Response. In preparation.
- Karrison, T. G., Huo, D., and Chappell, R. (2003). A group sequential, response-adaptive design for randomized clinical trials. *Controlled Clinical Trials*, 24(5):506–522.
- Kiss, I. Z., Miller, J., and Simon, P. (2017). *Mathematics of Epidemics on Networks From Exact to Approximate Models*. Springer.
- Klepac, P., Kucharski, A. J., Conlan, A. J., Kissler, S., Tang, M., Fry, H., and Gog, J. R. (2020). Contacts in context: large-scale setting-specific social mixing matrices from the BBC Pandemic project. *medRxiv*, page 2020.02.16.20023754.
- Kucharski, A. J., Klepac, P., Conlan, A., Kissler, S. M., Tang, M., Fry, H., Gog, J., and Edmunds, J. (2020). Effectiveness of isolation, testing, contact tracing and physical distancing on reducing transmission of SARS-CoV-2 in different settings. *medRxiv*, page 2020.04.23.20077024.
- Li, Q., Guan, X., Wu, P., Wang, X., Zhou, L., Tong, Y., Ren, R., Leung, K. S., Lau, E. H., Wong, J. Y., Xing, X., Xiang, N., Wu, Y., Li, C., Chen, Q., Li, D., Liu, T., Zhao, J., Liu, M., Tu, W., Chen, C., Jin, L., Yang, R., Wang, Q., Zhou, S., Wang, R., Liu, H., Luo, Y., Liu, Y., Shao, G., Li, H., Tao, Z., Yang, Y., Deng, Z., Liu, B., Ma, Z., Zhang, Y., Shi, G., Lam, T. T., Wu, J. T., Gao, G. F., Cowling, B. J., Yang, B., Leung, G. M., and Feng, Z. (2020). Early transmission dynamics in Wuhan, China, of novel coronavirus-infected pneumonia. *New England Journal of Medicine*, pages 1–9.
- Long, Q.-X., Liu, B.-Z., Deng, H.-J., Wu, G.-C., Deng, K., Chen, Y.-K., Liao, P., Qiu, J.-F., Lin, Y., Cai, X.-F., Wang, D.-Q., Hu, Y., Ren, J.-H., Tang, N., Xu, Y.-Y., Yu, L.-H., Mo, Z., Gong, F., Zhang, X.-L., Tian, W.-G., Hu, L., Zhang, X.-X., Xiang, J.-L., Du, H.-X., Liu, H.-W., Lang, C.-H., Luo, X.-H., Wu, S.-B., Cui, X.-P., Zhou, Z., Zhu, M.-M., Wang, J., Xue, C.-J., Li, X.-F., Wang, L., Li, Z.-J., Wang, K., Niu, C.-C., Yang, Q.-J., Tang, X.-J., Zhang, Y., Liu, X.-M., Li, J.-J., Zhang, D.-C., Zhang, F., Liu, P., Yuan, J., Li, Q., Hu, J.-L., Chen, J., and Huang, A.-L. (2020). Antibody responses to SARS-CoV-2 in patients with COVID-19. *Nature Medicine*, pages 1–4.
- Marzi, A., Robertson, S. J., Haddock, E., Feldmann, F., Hanley, P. W., Scott, D. P., Strong, J. E., Kobinger, G., Best, S. M., and Feldmann, H. (2015). VSV-EBOV rapidly protects macaques against infection with the 2014/15 Ebola virus outbreak strain. *Science*, 349(6249):739–742.
- Mehta, C., Gao, P., Bhatt, D. L., Harrington, R. A., Skerjanec, S., and Ware, J. H. (2009). Optimizing trial design. Sequential, adaptive, and enrichment strategies. *Circulation*, 119(4):597–605.
- Nason, M. (2016). Statistics and logistics: Design of Ebola vaccine trials in West Africa. *Clinical Trials*, 13(1):87–91.
- Proschan, M. and Evans, S. (2020). The temptation of response-adaptive randomization. *Infectious Diseases Society of America*, pages 1–8.

- Rosenberger, W. F., Stallard, N., Ivanova, A., Harper, C. N., and Ricks, M. L. (2001). Optimal Adaptive Designs for Binary Response Trials With Three Treatments. *Biometrics*, 57:909–913.
- Samai, M., Seward, J. F., Goldstein, S. T., Mahon, B. E., Lisk, D. R., Widdowson, M. A., Jalloh, M. I., Schrag, S. J., Idriss, A., Carter, R. J., Dawson, P., Kargbo, S. A., Leigh, B., Bawoh, M., Legardy-Williams, J., Deen, G., Carr, W., Callis, A., Lindblad, R., Russell, J. B., Petrie, C. R., Fombah, A. E., Kargbo, B., McDonald, W., Jarrett, O. D., Walker, R. E., Gargiullo, P., Bash-Taqi, D., Gibson, L., Fofanah, A. B., Schuchat, A., Neuzil, K., Insip, H., Risi, G., and Sow, S. (2018). The Sierra Leone Trial to Introduce a Vaccine Against Ebola: An Evaluation of rVSVΔG-ZEBOV-GP Vaccine Tolerability and Safety during the West Africa Ebola Outbreak. *Journal of Infectious Diseases*, 217(December 2014):S6–S15.
- Schoenfeld, D. A. . (1983). Sample-size formula for the proportional-hazards regression model. *International Biometric Society*, 39(2):499–503.
- Shim, E. and Galvani, A. P. (2012). Distinguishing vaccine efficacy and effectiveness. *Vaccine*, 30(47):6700–6705.
- Simon, R. and Simon, N. R. (2011). Using randomization tests to preserve type I error with response-adaptive and covariate-adaptive randomization. *Statistics Probability Letters*, 81(7):767–772.
- Tapiwa, G., Cécile, K., Dongxuan, C., Andrea, T., Christel, F., Jacco, W., and Niel, H. (2020). Estimating the generation interval for COVID-19 based on symptom onset data. *medRxiv*, pages 1–13.
- Teel, C., Park, T., and Sampson, A. R. (2015). EM estimation for finite mixture models with known mixture component size. *Commun Stat Simul Comput*, 44(6):1545–1556.
- Thall, P. F. and Wathen, J. K. (2007). Practical Bayesian adaptive randomization in clinical trials. *Eur J Cancer*, 43(5):859–866.
- Villar, S. S., Bowden, J., and Wason, J. (2015). Multi-armed bandit models for the optimal design of clinical trials: Benefits and challenges. *Statistical Science*, 30(2):199–215.
- Villar, S. S., Bowden, J., and Wason, J. (2018). Response-adaptive designs for binary responses: How to offer patient benefit while being robust to time trends? *Pharmaceutical Statistics*, 17(2):182–197.
- Wason, J. M. S., Brocklehurst, P., and Yap, C. (2019). When to keep it simple – adaptive designs are not always useful. *BMC Medicine*, 17(1):1–7.
- Weinberg, G. A. and Szilagyi, P. G. (2010). Vaccine epidemiology: Efficacy, effectiveness, and the translational research roadmap. *Journal of Infectious Diseases*, 201(11):1607–1610.
- WHO Disease Surveillance and Response Programme Area Disease Prevention and Control Cluster (2014). Contact tracing during an outbreak of Ebola virus disease. Technical Report September.
- WHO R&D Blueprint (2020). An international randomised trial of candidate vaccines against COVID-19. Technical report.
- Williamson, S. F., Jacko, P., Villar, S. S., and Jaki, T. (2017). Europe PMC Funders Group Europe PMC Funders Author Manuscripts A Bayesian adaptive design for clinical trials in rare diseases. pages 136–153.

A The COVID-19 model

There are two components to the model: the network that describes relationships between individuals who have the potential to be involved in the trial; and the transmission model that describes the dynamics of disease over the network.

The relationships in the network define the contact structure, and facilitate infection transmission. We make the assumption that our network for disease propagation includes (as a subnetwork) the network for contact tracing. The disease transmission tree is also a subnetwork of the whole network, and might include some edges not in the (relationship) contact network. In terms of modelling, these structures would need to change for a very large population: we consider here only populations of the order of one thousand people.

Both the network model and the disease transition model will be specific to a particular setting and, indeed, might change or become better informed over the course of the trial. Therefore, we use models without supposing that they will reflect the “truth” for any particular disease or scenario. Instead, we make choices we believe to be plausible in order to demonstrate the general methods, whose features will persist for settings such as these within the parameters we have described, i.e., diseases that pass from person to person, through contacts that can be traced.

A.1 Network model

We construct a network to represent relationships between individuals, where individuals are the nodes (or vertices) of the network, and relationships are edges between nodes. Infection is transmitted from one node to another across an edge that joins them. Therefore we want the network to include all possible transmissible relationships. These include interpersonal relationships, as well as random relationships, as it is likely that COVID-19 spreads between people who encounter each other only transiently as well as between people who know each other.

In our simulation, we consider a society in a moment where “social distancing” (or “physical distancing”) is insisted upon: schools are closed but employees go to work; social spaces such as hospitality and entertainment establishments are closed. Therefore we consider three types of relationships: within the home, in the workplace, and transient (or random). Transient relationships include encounters when in transit, or in shops, or due to the sharing of infrastructure. This is slightly different from the relationships considered in Kucharski et al. (2020), where interactions were localised to places (such as home or school or work) and within those spaces, a proportion of encounters are with people who are “acquaintances”, that is, people who have been met before, and a proportion of encounters will be with new people. In contrast, we designate that all work and home contacts are “acquaintances” and all random relationships are not, by definition.

For our simulation, we ascribe to individuals ages: below 19, 19 to 65, and over 65. We make this distinction as there are data available to guide the construction of the residency part of the network. People live in dwellings according to the UK census (<https://www.ons.gov.uk/peoplepopulationandcommunity/housing/adhocs/008634ct08192011censushouseholdtypehouseholdsizeandageofusualresidentshouseholdsendlandandwales>). We choose to simulate 500 dwellings in total. People aged 19 to 65, and one fifth of people aged 65+, are connected to approximately 15 other people via a “workplace”. The number of people in the workplace is reflective of the likely number of people with whom an infrastructure is shared, rather than the number of colleagues. Together, the household edges and the workplace edges are the predictable edges that make up the contact network (the network of acquaintances, in the terminology of Kucharski et al. (2020)). Finally, we include approximately 10 random connections per person to any other person. These are unpredictable (random/transient) edges.

There are many things we have omitted from this example, which might be important to include in a particular application. Omissions include: institutional residences, including care homes, halls of residence, barracks and prisons; structured contacts between children (such as schools); structured social contacts; extended family contacts; neighbour contacts; social spaces. The “workplace” element in our simulation lacks nuances such as details of which people in the population commute to a regular workplace, and with whom this workplace is shared. Such details could be filled in using information from the census and the Labour Force Survey for the UK.

It is also worth noting that these network structures might change, e.g. in response to the progress of the epidemic. These changes will affect the extent to which traced contacts account for new observed cases. If the fraction is smaller than anticipated, the study will be underpowered, as fewer events than expected will be observed in the trial population. To maintain the trial’s power, the key is to establish who the index cases shared space with in the last week and who these contacts will likely share space with in the coming week. Hopefully, then, it will still be possible to capture individuals at imminent risk of exposure in a natural way.




Figure 2: Example contact network showing 80 people in a contact network. “Predictable” edges between housemates and colleagues are shown. Individuals are coloured by age group and clustered into households.

A.1.1 Definitions

Our network is an undirected graph,

$$\mathcal{G} = (\mathcal{V}, \mathcal{E})$$

with vertices, or nodes,

$$\mathcal{V} = \{i, i = 1, \dots, N_I\}$$

representing the N_I individuals, which we index with i , and edges

$$\mathcal{E} \subseteq \{\{i, j\} | (i, j) \in \mathcal{V}^2 \cap i \neq j\}$$

representing connections, or relationships, between individuals. Each node i has a set of attributes $\{a_1(i), a_2(i), \dots, a_{N_a}(i)\}$.

The network operates at the level of the individuals. Each individual belongs to a household \mathcal{H}_h and every individual in a household is connected to every other individual in the household, as in Fyles et al. (2020), such that the induced subgraph $\mathcal{G}_1[\mathcal{H}_h]$ is completely connected:

$$\mathcal{E}_1[\mathcal{H}_h] = \{\{i, j\} | (i, j) \in \mathcal{V}[\mathcal{H}_h]^2 \cap i \neq j\}.$$

Let $\mathcal{H} = \{\mathcal{H}_h, h = 1, \dots, N_H\}$ be the set of N_H households, which forms a partition of the vertices \mathcal{V} (all subsets are mutually disjoint and their union is equal to the set). Let $a_1(i)$ be the household of individual i , i.e. $a_1(i) = h \Leftrightarrow i \in \mathcal{V}[\mathcal{H}_h]$. Then the size of a household is

$$\mathcal{O}(\mathcal{H}_h) = \sum_{i=1}^{N_I} \mathbf{1}_{a_1(i)=h}.$$

We start with $N_H = 500$ households with $\mathcal{O}(\mathcal{H}_h)$ people in each. The age and number distribution follows <https://www.ons.gov.uk/peoplepopulationandcommunity/housing/adhocs/008634ct08192011censushouseholdtypehouseholds>. That is, each person has an age attribute,

$$a_2(i) \in \{1, 2, 3\}$$

where $a_2(i) = 1$ if person i is under 19, $a_2(i) = 2$ if person i is 19 to 65, and $a_2(i) = 3$ if person i is over 65.

We define an individual i to be part of the workforce via a “worker” variable $a_3(i)$, where

$$\begin{aligned} a_2(i) = 1 &\implies a_3(i) = 0, \\ a_2(i) = 2 &\implies a_3(i) = 1, \\ a_2(i) = 3 &\implies \begin{cases} a_3(i) = 0 & \text{with probability 0.8} \\ a_3(i) = 1 & \text{with probability 0.2.} \end{cases} \end{aligned}$$

There are $A_3 = \sum_i \mathbf{1}_{a_3(i)=1} \approx 731$ people in the workforce. We define $A_3/15 \approx 49$ workplaces to which people with $a_3(i) = 1$ (that is, all those aged 19 to 65, and 1/5 of people over 65) are assigned following a multinomial distribution. We choose to place on average 15 people in a “workplace”, \mathcal{W}_w , which represents not their employment structure but a close shared use of the infrastructure. A useful guide might be how many toilets per person a place of work should have.

The workplaces are completely connected:

$$\mathcal{E}[\mathcal{W}_w] = \{\{i, j\} | (i, j) \in \mathcal{V}[\mathcal{W}_w]^2 \cap i \neq j\}.$$

Finally, 1000 totally random edges are added, which amounts to around ten per person:

$$\Pr(\{i, j\} \in \mathcal{E}) = \frac{10}{(\sum_h \mathcal{O}(\mathcal{H}_h))}.$$

These edges correspond to potential transmission encounters that would not be recalled or anticipated through contact tracing. The result is an average of 20 connections per person, of which ten have a weight of 1 and ten have a weight of 0.1.

A.1.2 Parametrisation

The parameters we have used and their provenance are listed in Table 7.

Table 7: Parameters used in the COVID-19 disease transmission and vaccine trial model. $\mathcal{N}(l, \mu, \sigma)$ denotes a normal distribution truncated at l .

Parameter	Value	Source
Number of households	500	Chosen with reference to other choices
Household size	Draws from raw data	UK 2011 census
Workplace size	$\sim \text{Poisson}(15)$	https://www.hse.gov.uk/contact/faqs/toilets.htm
β (per-contact infection or transmission rate)	0.01	Chosen with reference to other choices
Predictable edge weight (χ_h, χ_w)	1	Chosen with reference to other choices
Unpredictable edge weight (χ_n)	0.1	Chosen with reference to other choices
Incubation period	$\sim 2 + \Gamma(\text{shape} = 13.3, \text{rate} = 4.16)$	Li et al. (2020)
Infectious period	$\sim 1 + \Gamma(\text{shape} = 1.43, \text{rate} = 0.549)$	Li et al. (2020)
Time to enrol whole contact network	$\sim \mathcal{N}(l = 0, 10.32, 4.79)$	From Henao-Restrepo et al. (2017). One could use ECDC (2020). Fyles et al. (2020) use Poisson Parameter $\sim \text{Uniform}(1.5, 2.5)$
Time to vaccine-induced sero-conversion	$\sim \Gamma(\text{shape} = 3, \text{rate} = 1)$	Plausibly all seroconverted within 14 days

A.2 Disease and trial state transitions

Individuals' disease states and possible transitions are described by a compartmental model (Figure 3), as in Camacho et al. (2015). Possible state transitions are disease progression ($S \rightarrow E \rightarrow I_1 \rightarrow I_2 \rightarrow R$, where S denotes Susceptible, E exposed, I infectious, and R removed), similarly to Danon et al. (2020), and trial enrolment. There are rules for both types of transitions. Transition rates depend on the rules and the states of the neighbours.




Figure 3: Disease-state transition model for members of the population who aren't enrolled (U) and those enrolled and vaccinated (V) and those enrolled to the control group (C). Arrows show possible transitions between states, labelled by the rates. Everyone starts in State S_U , except for one person who starts in E_U . Rates of transition of person i depend on (a) the current state of person i and (b) the current states of the neighbourhood of person i . The instantaneous rate of infection (or infection hazard), k_x , depends on vaccination status x . People spend ξ days in the exposed state before moving to the infectious state. In the first phase of infectiousness, I_1 , all people are infectious and asymptomatic for one day (Kucharski et al., 2020). 20% of people remain asymptomatic for a further infectious period of γ days. The other 80% are symptomatic for γ days.

Our simulated trial operates on a time unit of one day, and has total duration of the order of 100 days. For simplicity, we assume that we begin initiation of one contact network (corresponding to one index case) every day. Initiation is the moment where all nodes in the network are in state S_U except one who is on their first day as state E_U . Enrolment begins when this node reaches state I_2 .

Each participant who is enrolled has as their reference day the day on which they were enrolled. Therefore every participant has their own individual timeline which can be referenced to others by the day of their enrolment.

A.2.1 Disease transmission rules

A person in the E state becomes infectious, I_1 , 24 hours before entering the I_2 state as in Kucharski et al. (2020). They are infectious but not symptomatic. For our simulation, we assume state/society rules that symptomatic persons are to shelter in place (which means, for our simulation, they do not leave their homes). Therefore, as an infectious I_1 , a person can infect their home contacts, their work contacts and their random contacts. The majority of I_1 people become I_2 people. As an I_2 , a person can infect only their home contacts. Some infectious people, however, are not symptomatic, so these people behave as I_1 s, in terms of who they infect. This is represented by an additional edge from I_1 to R , bypassing I_2 , where the person stays in the I_1 state for an addition period, equal to the period they would have spent as an I_2 if symptomatic. Transitions to state R imply removal from the infectious population: this can be due to recovery or hospitalisation, and hence isolation.

A susceptible individual i belonging to arm x becomes “exposed” (i.e. transitions to state E) with rate

$$k_x(i) = \beta \left(\chi_h \mathcal{M}_i^{(h)} + \chi_h \mathcal{L}_i^{(h)} + \chi_w \mathcal{L}_i^{(w)} + \chi_n \mathcal{L}_i^{(n)} \right) (1 - \eta \cdot x_i), \quad x \in \{V, U, C\},$$

where h refers to household, w to workplace, and n to random. $\chi_h > 0$ is the scalar (or edge weight) for household contacts, $\chi_w > 0$ is the scalar (or edge weight) for workplace contacts, and $\chi_n > 0$ is the scalar (or edge weight) for random contacts. One could specify a different χ_h for pre- and post-symptomatic infectiousness, or even a changing profile over time, as described in He et al. (2020). $x_i = 1$ if person i is vaccinated and 0 otherwise. $\eta \leq 1$ is the vaccine efficacy, and $\beta > 0$ is the per-contact rate at which infectious people infect their susceptible contacts.² $\mathcal{M}_i^{(h)}$ is the number of I_2 household (h) contacts of susceptible i .

$$\mathcal{L}_i^{(y)} = \sum_{j \in \mathcal{N}_{(y)}(i)} \mathbf{1}_{j \in \{I_1, \cdot\}}, \quad y \in \{h, w, n\}$$

²The unit of β is per contact per day; the scalars χ are unitless. Thought of as edge weights, they transform the static (or quenched) network into a dynamic or adaptive network through link deactivation (Kiss et al., 2017).

is similarly defined as the number of contacts who are infectious but not symptomatic.

Note that there is no infection other than from a contact (with the exception of the simulated “time trends”, Figure 5). Thus any infection within a contact network came (directly or indirectly) from the contact network’s index case.

Our model yields new infection events as occurring from pre-symptomatic people 58% of the time and from symptomatic people 42% of the time, omitting transmissions from people who never become symptomatic. In comparison to reported estimates (44% from 77 recorded transmission pairs (He et al., 2020), and 48 and 62% in Singapore and Tianjin, respectively (Tapiwa et al., 2020)), we confirm that our simulation scenario is consistent with one with quick quarantine of close contacts (He et al., 2020).

A.2.2 Trial rules

Our simulation mimics the trial rules of the ring recruitment strategy (Henao-Restrepo et al., 2015), with the exception that we randomise at the individual, rather than cluster, level, and therefore do not exclude individuals on the basis of contact-network overlap with existing enrolled networks.

In our trial we define contact tracing as identifying only existing relationships (“acquaintances”, or “those met before”, in the terminology of Kucharski et al. (2020)) that might be, or might have been, a means for transmission. That is, a newly diagnosed person is asked to recall all of their contacts, and these individuals are contacted and asked likewise to list all their contacts. This makes our simulation and trial design most like the “self-isolation and manual contact tracing of acquaintances” of Kucharski et al. (2020). The relationships that are of interest will depend on the society and any concurrent actions, guidance or instruction from the state, which in our simulation include home and workplace contacts, while random contacts remain unpredictable. We assume unpredictable relationships are not recalled in contact tracing and, similarly, cannot be anticipated for recruitment of trial participants. We assume work and home relationships are recalled perfectly. By recruiting “contacts” and “contacts of contacts” into the trial (the “two-step tracing” of Fyles et al. (2020)), we are recruiting the people the index case lives with and their colleagues, and those the case works with and their housemates.

For our purposes, an index case for a contact network is a person identified as being in state $I_{2,}$ after the initiation of the trial. Eligible people are traced as described in ECDC (2020) and those who consent are enrolled as soon as they are identified and give their consent. Susceptible and exposed people are eligible for enrolment if they are a “predictable” contact of the index case and they are not already enrolled in the trial. Symptomatic $I_{2,U}$ people are excluded on the basis of their symptoms. R_U are excluded on the basis of their history, which we assume a perfect knowledge of in our simulation. This could result either from people being able to identify having had COVID-19, or from there being an accurate and reliable antibody test. Inclusion of R_U people in the simulated trial would result in a dilution of infections and would therefore require enrolment of more participants to maintain power. Their inclusion is advocated for reasons of safety testing (<https://www.who.int/publications/i/item/an-international-randomised-trial-of-candidate-vaccines-against-covid-19>, <https://clinicaltrials.gov/ct2/show/NCT04405076>). Elsewhere seropositive people are excluded from vaccine trials (<https://www.clinicaltrialsregister.eu/ctr-search/trial/2020-001228-32/GB>).

The enrolment rate is $0 \leq \epsilon \leq 1$, which is the probability for each eligible person to enrol, where we use only contact structure and lack of symptoms to define eligibility. We use $\epsilon = 0.7$.³ Enrolled participants are randomised to the experimental arm with probability $0 \leq \pi_1 \leq 1$, and $\pi_1 + \pi_0 = 1$. Enrolment takes time δ . Those vaccinated have an additional wait time before reaching state S_V , which is development of immunity (or time to seroconversion), and which takes time τ . Transition from S_U to E_U is possible in this wait time. We assume that the vaccine effect before seroconversion is zero and that the vaccine effect after is the full effect of the vaccine.

Individuals who are in state E and are unenrolled are enrolled with the same probability (ϵ) as the susceptibles S_U , as they are asymptomatic. The same wait time for enrolment applies, but time to seroconversion does not, as the individual is already infected. If the participant transitions to $I_{2,U}$ before the recruitment time elapses, they will be excluded from the trial, as they will be showing symptoms and can be tested for confirmation.

Result accrual relies on surveillance and self reporting. For our simulations we assume that a fraction 0.2 of infectious individuals are asymptomatic. These infection events go unreported in the trial results (but still contribute to onward transmission).

We simulate one contact network at a time, beginning when the index case is identified. Transmission

³Henao-Restrepo et al. (2017) report that 50% of people identified were eligible and enrolled, where their eligibility criteria excluded people who were pregnant, breastfeeding, or under the age of 18.

in each contact network is independent of all other contact networks in the trial. The contact networks are related only through the time reference, in that one contact network is initiated on each day. Ideally, one would simulate a whole population that encompasses potentially thousands of distinct contact networks, so that each contact network might have a history that is affected by the events stemming from other contacts, and the moments for contact-network enrolment would arise in a more natural way. Here, however, we present only an approximation to that population, and omit entirely any shared history, save the accrued results that determine the response-adaptive allocation probabilities in the adaptive designs.

Finally, note that, for simplicity, we assume that randomisation, enrolment and vaccination are all assumed to happen on a single day for each individual, although the day will differ between individuals.

B Analyses, and exclusion criterion implemented at analysis points

Here we detail all equations to accompany Section 3.2, which describes the different ways we explore to analyse the outcome. We present the methods in the same order and use a single framework that describes all the methods in the same way.

B.1 Analysis of raw data

We have $j = 1, \dots, N_I$ individuals. Each has a vaccination status, x_j , and a disease status y_j , where the vaccination status is dictated by the trial design and the disease status from the underlying epidemic model:

$$x_j = \begin{cases} 0 & \text{person } j \text{ not vaccinated} \\ 1 & \text{person } j \text{ vaccinated} \end{cases}$$

and

$$y_j = \begin{cases} 0 & \text{person } j \text{ not diagnosed} \\ 1 & \text{person } j \text{ diagnosed} \end{cases}$$

at the end of the trial.

The test statistic is a standard normal variable Z , where

$$Z = \frac{\hat{p}_1 - \hat{p}_0}{\sqrt{\sigma_0 + \sigma_1}}, \quad (1)$$

$$\hat{p}_v = \frac{\sum_{j:x_j=v, y_j=0} \omega_j}{N_v}, \quad (2)$$

$$N_v = \sum_{j:x_j=v} \omega_j \quad (3)$$

$$\omega_j = 1 \quad \forall j,$$

and

$$\sigma_v = \frac{\hat{p}_v(1 - \hat{p}_v)}{N_v}. \quad (4)$$

p_v is the true probability of not being a confirmed case if in arm v , and \hat{p}_v is our estimate of it, defined as the proportion of people in arm v not confirmed, where $v = 0$ is the control arm and $v = 1$ is the experimental arm; N_v is the total number in arm v ; and σ_v is the variance of the estimator \hat{p}_v . Power is defined as the proportion of Z values that exceed 1.64, which is the 95th quantile of a standard normal distribution. ω_j is the weight of person j which, for the unweighted method, is 1 for all participants. In the descriptions that follow, we see that the weights ω define the retrospective exclusion criterion so that all methods use the same calculation and each is defined only by the definition of the weights.

B.2 Analysis using binary weighting

The binary-weighting method proceeds as above but considers also the day of commencement of symptoms, s_j . Individuals are excluded if $s_j < 9$ relative to a randomisation day of 0. We write this as the weight, ω_j , for each individual j , so that a weight of 0 equates to exclusion:

$$\omega_j = \begin{cases} 0 & s_j < 9 \\ 1 & s_j \geq 9 \text{ or } y_j = 0 \end{cases}$$

These are used together with Equations 1, 2, 3, and 4 as before.

B.3 Analysis using continuous weighting

Given person j 's symptoms began on day s_j relative to their randomisation day of 0, τ and ξ are their unknown time to seroconversion⁴ and incubation time, respectively. The probability they were infected after the trial began is $P(\tau + \xi < s_j)$. We assume τ and ξ are distributed between individuals as $\Gamma(\text{shape} = 3, \text{rate} = 1)$ and $2 + \Gamma(\text{shape} = 13.3, \text{rate} = 4.16)$. Hence the distribution of $\tau + \xi$ is estimated by matching moments using the Welch-Satterthwaite equation, as described in Box (1954).

⁴Note that we include time to seroconversion also for the control group, who don't receive the vaccine, and don't seroconvert.

We estimate the vaccine efficacy $0 \leq \hat{\eta} \leq 1$ as

$$\hat{\eta} = 1 - \frac{f_1}{N_1} \bigg/ \frac{f_0}{N_0},$$

where

$$\hat{f}_v = \sum_{j: x_j=v, y_j=1} \omega_j.$$

Suppose from the cumulative distribution function of the gamma distribution we have a nominal probability, i.e. neglecting the effect of the vaccine, q_j that the day person j was infected, d_j , was after seroconversion on day D_c . We write the complement, the probability that person j was infected before day D_c , as $r_j = 1 - q_j$. We re-estimate their probability given that person j was vaccinated ($x_j = 1$):

$$\Pr(d_j > D_c | x_j = 1) = \frac{(1 - \hat{\eta})x_j}{y_j + (1 - \hat{\eta})x_j}, \quad (5)$$

because, if there is some efficacy, then they are more likely to have been infected before being vaccinated than after (relative to a vaccine that has no effect: $\Pr(d_j > D_c | x_j = 0) = q_j$). We solve this iteratively for $\hat{\eta}$ with reference to all observations j .⁵ Then

$$\omega_j = \begin{cases} 1 & y_j = 0 \\ \Pr(d_j > D_c | x_j) & y_j = 1 \end{cases} \quad (6)$$

The resulting weights are used in Equations 1, 2, 3, and 4 as before.

Determining the inclusion weight for a vaccinated person

The inclusion weight for a vaccinated person j ($q_j = 1$), whose symptoms began after D_c (randomisation date plus an assumed time from vaccination to seroconversion) is the probability that they were infected after D_c . To determine this, first we make explicit the condition that their infection date d_j is less than their symptom date D_s . The conditional probability we want can then be decomposed into probabilities unconditional on the symptom date, as follows,

$$P(d_j > D_c | q_j = 1, d_j < D_s) = P(D_c < d_j < D_s, q_j = 1) / P(d_j < D_s, q_j = 1)$$

Then splitting the denominator into the probabilities of being infected in two different periods (before versus after D_c) gives

$$P(d_j > D_c | q_j = 1, d_j < D_s) = \frac{P(D_c < d_j < D_s | q_j = 1)}{P(d_j < D_c | q_j = 1) + P(D_c < d_j < D_s | q_j = 1)}$$

Since the probability of being infected before D_c doesn't depend on whether person j was vaccinated, $P(d_j < D_c | q_j = 1) = P(d_j < D_c | q_j = 0)$. We additionally assume that

$$P(D_c < d_j < D_s | q_j = 1) = \psi P(D_c < d_j < D_s | q_j = 0)$$

where ψ is the relative risk of infection between a vaccinated and unvaccinated person, assumed to be constant through time, giving

$$P(d_j > D_c | q_j = 1, d_j < D_s) = \frac{\psi P(D_c < d_j < D_s | q_j = 0)}{P(d_j < D_c | q_j = 0) + \psi P(D_c < d_j < D_s | q_j = 0)}$$

Expressing the right hand side in terms of probabilities conditional on symptoms, by dividing the numerator and denominator by $P(d_j < D_s | q_j = 0)$, then gives the weight for a vaccinated person j as

$$\omega_j = P(d_j > D_c | q_j = 1, d_j < D_s) = \frac{\psi \omega_j^{(0)}}{(1 - \omega_j^{(0)}) + \psi \omega_j^{(0)}}$$

where $\omega_j^{(0)} = P(D_c < d_j < D_s | d_j < D_s, q_j = 0) = P(D_c < d_j < D_s | q_j = 0) / P(d_j < D_s | q_j = 0)$ is the weight if person j were unvaccinated.

⁵This is solved as in expectation maximisation (EM). EM algorithms have been used and discussed in clinical trials for e.g. sample-size re-estimation (Gould and Shih, 1992; Friede and Kieser, 2002; Teel et al., 2015; Huang et al., 2018).

C Supplementary material

C.1 Weighting example

Table 8: Analysis methods for trials simulated under the “binary” method framework corresponding to Table 3. We report the total weights for vaccinated and control confirmed cases. The row “True at recruitment” is the result we would obtain with perfect knowledge of who was infected on the day of recruitment. “True” is the result we would obtain if we had perfect knowledge of the day of infection and the day of seroconversion.

Method	Power	Type 1 error	Vaccinated	Control	VE estimate
None	0.59	0.07	16.5	27.8	0.39 (0.19)
Binary	0.74	0.04	6.9	18.0	0.59 (0.19)
Continuous (without VE)	0.75	0.03	6.6	17.0	0.59 (0.18)
Continuous (with VE)	0.81	0.04	5.7	17.0	0.64 (0.19)
True at recruitment	0.72	0.04	8.5	19.7	0.55 (0.18)
True	0.93	0.09	5.2	19.7	0.72 (0.14)

C.2 Allocation probabilities




Figure 4: Five samples of trajectories of the allocation probability over time for the adaptive designs.

C.3 Time trend in incidence of infection

Figure 5 shows the robustness of one adaptive trial design to a time trend, demonstrating the effect of the correction of Simon and Simon (2011). We choose to illustrate with linear trends rather than something more realistic, undulating, or stochastic in order to stress test the method. In addition, we choose a trend that we would expect to most favour the Thompson sampling methods. Indeed, we find that uncorrected Thompson sampling sees an increase in type 1 error, which is corrected by the resampling method of Simon and Simon (2011). As TS is a more “aggressive” response-adaptive method than TST, we expect it to be more susceptible to time trends, and therefore the loss in power following correction to be exaggerated.




Figure 5: The robustness of one adaptive trial design (TST) to a time trend, demonstrating the effect of the correction of Simon and Simon (2011). Type 1 error rate and power as a function of the trend in the background rate for a trial with a response-adaptive randomisation rate. “Background rate” can be interpreted as the rate of infection by individuals unknown in the context of the trial. (E.g. source population or unknown contact.) We illustrate with linear trends rather than something more realistic, undulating, or stochastic in order to stress test the method. In addition, we choose a trend that we would expect to most favour the Thompson sampling methods. As TS is a more “aggressive” response-adaptive method than TST, we expect it to be more susceptible to time trends, and therefore the loss in power following correction to be exaggerated. Left: Five different time trends for the background rate. The trend is that background rates diminish over time to zero. Middle: The gradient of the trend is shown on the x axis. On the y axis is the type 1 error rate. Right: The gradient of the trend is shown on the x axis. On the y axis is the power.

D Trial size

We compare our fixed-case-weight to a fixed-sample-size design in terms of the power: do different possible epidemic realisations result in predictable or uncertain powers for the trial designs under the different rules governing the final trial size?

To assess the utility of defining trial size according to cumulative observed cases, we simulate three scenarios: one, where the parameters for the simulated trials are the same as those for the trial design, and ones where the rate of transmission for the simulation is $4/5$ and $5/4$ what was used for the design. Varying the rate of transmission serves to deviate the trajectory of number of participants and cases observed systematically from what is expected at the “Design” stage when you must decide a rule to stop the trial before observing any data. This allows us to show how noise in the infection model affects the power of trials designed with a fixed number of participants, compared to trials designed with a fixed number of cases. We additionally compare the two designs where there is no vaccine efficacy, demonstrating how many resources are consumed in order to conclude futility.

Note that this artificial discretisation of the outcome space into three represents the variation we might see due to stochasticity, or due the model systematically under or overestimating the number of cases. Indeed, in a trial in an unfolding epidemic, one would likely not be able to tell whether, or to what extent, the deviance is due to stochasticity and to what extent it is due to the model being systematically wrong. Therefore, we use different model parametrisations to investigate the effects of the different epidemic realisations on power subject to the design of the trial, and note that we observe the same effect even when the model of the data-generating mechanism is “correct” (Table 10).

In Table 9, we compare a design where we terminate the trial once we have recruited a fixed number of contact networks (60) to one that terminates once we have observed a minimal number (24) of events. In our “base case”, using the parameters described in the Appendix, the two approaches have the same expected power (0.8). The expected number of cases for 60 contact networks is 24, and the expected number of contact networks for 24 cases is 60. Differences emerge when we compare the case where the vaccine has no effect, and where the outcome diverges from the expected values.

Where there is no vaccine effect, the fixed-number-of-participants method enrolls ≈ 800 more participants than the fixed-case method on average. Similarly, if we observe more events per participant than we expected under the experimental, the number of participants for the fixed-number-of-participants method enrolls ≈ 700 more participants than the fixed-case method on average. Conversely, if we see fewer events than we expected, the fixed-number-of-participants method concludes underpowered with 1,939 participants, and the fixed-case method enrolls 1,000 more participants but maintains power.

This presents another trade-off, as described by Mehta et al. (2009): if the trial size is defined by a pre-determined number of participants, we observe a range in numbers of cases, and a range in power. If the trial size is defined by a pre-determined number of cases, we observe a range in numbers of participants, but the power remains constant. Put another way, fixing sample size gives rise to uncertainty in the power, and fixing case numbers gives rise to uncertainty in the sample size, but the power can be expected to remain around 0.8, whatever happens in the epidemic.

Table 9: Comparison of designs where the trial size is determined by the number of cases to designs where the trial end is determined by the number of people enrolled. We make comparisons under three scenarios: one, where we assume that the trajectory we estimated was correct on average (“Average”), and ones where it is 4/5 on average (“Lower”) and 5/4 on average (“Higher”) what we estimated. Participants are recruited following the ring strategy. The trial follows the FR design with a follow-up time of 25 days and weights exclusion with a continuous variable that accounts for VE. Standard deviations in brackets.

Trajectory	Trial size determined by...	Number of participants	Duration (days)	Number of confirmed cases	Vaccinated	Power	VE estimate	Prevented export infections	Number of participants (no vaccine effect)	Type 1 error
Average	Cases 24–26	1920 (540)	85 (17)	52 (8)	960	0.80	0.65 (0.18)	4.99	1192 (341)	0.04
Lower	Cases 24–26	3017 (793)	119 (25)	57 (10)	1508	0.79	0.64 (0.2)	2.34	1836 (548)	0.04
Higher	Cases 24–26	1235 (332)	63 (10)	48 (7)	617	0.81	0.64 (0.18)	9.88	894 (144)	0.04
Average	Number of participants 1700–1842	1939 (115)	85 (4)	52 (12)	970	0.81	0.65 (0.19)	4.94	1939 (115)	0.05
Lower	Number of participants 1700–1842	1939 (115)	85 (4)	37 (9)	970	0.62	0.63 (0.28)	1.97	1939 (115)	0.06
Higher	Number of participants 1700–1842	1939 (115)	85 (4)	76 (16)	970	0.91	0.65 (0.14)	9.68	1939 (115)	0.05

We demonstrated the trade-off between defining the size of the trial based on the effective number of cases (resulting in uncertainty in number of participants), and on the effective number of participants enrolled (resulting in uncertainty in power). We advocate choosing the former, as the benefits of being able to anticipate the power exceed the benefits of being able to anticipate the number of participants. Indeed, in outbreaks, the number of participants is often uncertain without it being in the design. We note that the suggestion to end a trial based on the number of cases is proposed also in WHO R&D Blueprint (2020), suggesting that uncertainty in sample size is preferable to uncertainty in power in the context of vaccine trial in an outbreak.

Where there is no vaccine effect, defining trial size based on the total events seen results in a smaller number of participants, and therefore the trial can conclude more quickly, using fewer resources. Where there is a vaccine effect, defining trial size according to number of participants results in more uncertainty in power. The uncertainty in power arises as a result of uncertainty in the number of cases. In some circumstances, it might be that using the number of participants or both the number of cases and the number of participants will be the best way to control power. Because of the dependence on epidemic, disease and vaccination dynamics, the rule to choose, and the components of the rule, become apparent only through simulation.

On the x axis is N , the number of participants, and on the y axis is f , the weighted number of cases. Then π for a single point is the probability to reject the null hypothesis H_0 :

$$\pi(N, f) = \int_{f_0=0}^f p(f_0|N, f, \hat{\eta}) \cdot (Z(N, f_0, f) > C_0) df_0 \quad (7)$$

$$Z(N, f_0, f) = \frac{p_1 - p_0}{\sqrt{\sigma_0 + \sigma_1}} \quad (8)$$

$$p_v = 1 - f_v/(N/2) \quad (9)$$

$$f_1 = f - f_0 \quad (10)$$

$$\sigma_v = p_v(1 - p_v) \quad (11)$$

For this we assume the design D has fixed and equal randomisation. Then each arm has $N/2$ participants, and we integrate over all possible $\{f_0, f_1\}$ such that $f = f_0 + f_1$ and given $\hat{\eta}$.

I would define π the power of a statistical test associated with a point in (N, f) space, and $\hat{\pi}$ the power of a trial design, which is the expected power over the points in the design space. I've used the following definition for "power": "the statistical power (π) is then defined as the probability of rejecting the null hypothesis when it is false following some statistical test."

We write the power for a design as $\hat{\pi}$. A design is defined by a single N^* (or f^*), which is the point at which the trial stops. To compute $\hat{\pi}(N^*)$ (or $\hat{\pi}(f^*)$), we integrate (or sum) over f (or N), weighted by the probability to have observed f given N^* , $p(f|N^*, M, D)$ (or N given f^* , $p(N|f^*, M, D)$) which depends additionally on the model M :

$$\hat{\pi}(N^*) = \sum_f \pi(N^*, f) p(f|N^*, M, D),$$

and

$$\hat{\pi}(f^*) = \sum_N \pi(N, f^*) p(N|f^*, M, D).$$

That is, $\hat{\pi}$ is an expectation of π over f (or N). The variance in the f dimension of π is averaged over to compute $\hat{\pi}(N^*)$, the expected power for a trial that stops according to N^* . The variance in the N dimension of π is averaged over to compute $\hat{\pi}(f^*)$, the expected power for a trial that stops according to f^* . Figure 6 shows a greater range in values in π across the f dimension than the N dimension, suggesting higher variance in $\pi(N^*, f)$ than $\pi(N, f^*)$ for any N^* or f^* . Indeed, $\pi(N, f^*)$ is more-or-less constant for many values of f^* , so the variance in power for a design set to terminate at f^* has variance close to 0. There is higher variance in the π , then, for a trial ending at N^* than a trial ending at f^* . Thus the fixed-number-of-participants design has a higher probability of an underpowered study and a higher probability of recruiting more participants than necessary.

A trial ending by case number is less reliant on modelling the data-generating mechanism than a trial ending by participant number. At its limit, if we are able or willing to assume that $\pi(N, f^*)$ is constant for many values of f^* , we can write $\pi(N, f^*) = \pi^*(f^*)$ where $\pi^*(f^*)$ is a function only of f^* . Revisiting the equation for the

power for trial ending at f^* cases, and using also the property $\sum_N p(N|f^*, M, D) = 1$:

$$\hat{\pi}(f^*) = \sum_N \pi(N, f^*) p(N|f^*, M, D) \quad (12)$$

$$= \sum_N \pi^*(f^*) p(N|f^*, M, D) \quad (13)$$

$$= \pi^*(f^*) \sum_N p(N|f^*, M, D) \quad (14)$$

$$= \pi^*(f^*), \quad (15)$$

and thus the power for trial ending at f^* cases would not depend at all on the data-generating mechanism, our assumed model M . However, we would want to use simulation to check the relationship $\pi(N, f^*) = \pi^*(f^*)$, in particular the range of (N, f) values for which it is valid, and where the randomisation ratio is not fixed to be equal. We expect the relationship to weaken where incubation time or time to seroconversion is long.




Figure 6: Power, shown by colour (left) and boundary lines (right), to detect vaccine efficacy of 0.7, with results from Table 9 written in, showing three examples of power having terminated the trial after a prespecified number of participants and after a prespecified case weight. On the x axis is the cumulative weight of enrolled participants, a value less than or equal to the total number of people enrolled. On the y axis is the cumulative weight of all participants who are confirmed infected. Again, the effective number of cases will be less than the total number of people confirmed infected. Where all participants have a weight of 1, the total weight is equal to the number of participants and the total weighted cases is the number of cases. Where some participants are downweighted, these serve as upper bounds. The trial in real time will unfold as a monotonically increasing trajectory starting at $(0,0)$, in the bottom-left-hand corner. Each point on this trajectory, defined by the cumulative participant weight and effective number of cases, will have a characteristic power: that is, the expected fraction of null hypotheses correctly rejected.

The colours cover the spread of realisations from the standard model. The numbers correspond to three different trajectories from different models: one of increased incidence, the standard model, and reduced incidence. This shows that the biased models we consider, that differ systematically from the norm, have expected values still consistent with what we expect to see from the standard model. E.g. for the standard model, the probability to land on an outcome yielding power between 0.79 and 0.85 is 0.11 (corresponding to the middle “0.8”).

Table 10: The composition of power in two trial designs, corresponding to Figure 6 and analogous to Table 9. We make comparisons under three scenarios: one, where we observe the case incidence (trajectory) we expected, one where it is lower, and one where it is higher, by artificially discretising the result space into three across each axis. Each scenario has a probability to occur which together sum to one. Each scenario has a power, so that the power of the trial is the sum of the scenario powers weighted by their probability to occur: the expected power of the trial whose size is determined by the number of cases is $0.82 \cdot 0.1 + 0.83 \cdot 0.5 + 0.81 \cdot 0.4$, and that for a trials whose size is determined by the number of participants is $0.82 \cdot 0.1 + 0.74 \cdot 0.4 + 0.89 \cdot 0.5$. Notice how the former averages over three similar powers and the later over a range of powers. This means that the “Probability to occur” – which is what we use simulation for – is much more important for the latter case.

Trajectory	Trial size determined by...	Subset	Number of participants	Case weight	Power	Probability to occur
Average	Cases 24–26	Number of participants 1700–1842	1769	25	0.82	0.1
Lower	Cases 24–26	Number of participants >1842	2253	25	0.83	0.5
Higher	Cases 24–26	Number of participants <1700	1356	25	0.81	0.4
Average	Number of participants 1700–1842	Cases 24–26	1769	25	0.82	0.1
Lower	Number of participants 1700–1842	Cases <24	1773	19	0.74	0.4
Higher	Number of participants 1700–1842	Cases >26	1768	33	0.89	0.5

E Glossary by letter

Letter	Meaning	Letter	Meaning	Letter	Meaning	Letter	Meaning
α	Significance threshold	a	node attributes	A			
β	per-contact infection or transmission rate	b		B			
γ	infectious period	c		C	control-arm label		
δ	recruitment time	d	incubation time	D			
ϵ	enrolment rate	e	trial end	E	exposed	\mathcal{E}	set of edges
ζ		f	fails	F			
η	vaccine efficacy	g		G		\mathcal{G}	graph
θ		h	“household contact” label	H		\mathcal{H}	household
ι		i	index over node	I	infectious		
		j	index over node	J			
κ		k	instantaneous rate of infection	K			
λ		l	Truncated normal distribution parameter	L		\mathcal{L}	number infectious contacts
μ		m		M		\mathcal{M}	number symptomatic contacts
ν		n	“unpredictable contact” label	N	total	\mathcal{N}	neighbourhood
ξ	incubation period	o		O		\mathcal{O}	size
π	allocation probability	p	efficacy probability	P			
		q		Q			
ρ	allocation ratio	r		R	removed		
σ		s	day of symptoms	S	susceptible		
τ	seroconversion time	t	exposure time	T			
v		u		U	unenrolled label		
ϕ	Thompson tuning parameter	v	trial arm	V	experimental-arm label	\mathcal{V}	set of nodes
		w	“workplace contact” label	W		\mathcal{W}	workplace
χ	edge weight	x	vaccination label	X			
ψ	relative risk of infection given vaccination	y	confirmed-case label	Y			
ω	inclusion weight	z		Z			